

Announcements

- Azure credits have been issued (\$50)
 - Sent to @uw.edu emails
 - Post on Piazza if you have issues accepting it
 - Will use for HW3



Accept your Azure lab assignment

You have a pending lab assignment. Please accept your assignment to get started with your course.

[Accept lab assignment >](#)

This email is generated from an unmonitored alias; please do not reply. If you have questions, please [submit a request](#).

Recap: Grouping

```
SELECT Product, SUM(quantity)
FROM Purchases
GROUP BY Product
HAVING SUM(quantity) > 20
```

Product	Price	Quantity	Month
Bagel	3	20	Jan
Bagel	1.50	20	Feb
Banana	0.5	50	Feb
Banana	5	10	March
Apple	4	10	March



Product	SUM(quantity)
Bagel	40
Banana	60

Recap: Semantics

First evaluate the FROM clause

Next evaluate the WHERE clause

Group the attributes in the GROUPBY

Eliminate groups based on HAVING

Sort the results based on ORDER BY

Last evaluate the SELECT clause

FWGHOS™

Recap – General form

SELECT	S
FROM	R_1, \dots, R_n
WHERE	C1
GROUP BY	a_1, \dots, a_k
HAVING	C2
ORDER BY	O

S, O = any attributes a_1, \dots, a_k and/or any aggregates, but no other attributes

C1 = any condition on the attributes in R_1, \dots, R_n

C2 = any condition on the aggregate expressions and attributes a_1, \dots, a_k

Goals for Today

- We've completed our general form of a query
- Use SQL queries to assist other SQL queries
- Conclude our unit on SQL queries
 - After today you'll have essentially all the building blocks of most all queries you can think of

Outline

- Subquery mechanics
 - SELECT
 - FROM
 - WHERE/HAVING
- Decorrelation and unnesting along the way
- The Witness Problem

Nested Queries

A subquery is a SQL query nested inside a larger query

A subquery may occur in:

- A SELECT clause
- A FROM clause
- A WHERE or HAVING clause

Rule of thumb:

Avoid nested queries when possible...

...but sometimes it's impossible

Subqueries in SELECT

- Must return a single value
 - A 1x1 relation – single row, single column
- Uses:
 - Compute an associated value

Subqueries in SELECT

- Must return a single value
 - A 1x1 relation – single row, single column
- Uses:
 - Compute an associated value

```
SELECT P.Name, (SELECT AVG(P1.Salary)
                FROM Payroll AS P1
                WHERE P.Job = P1.Job)
FROM Payroll AS P
```

Subqueries in SELECT

- Must return a single value
 - A 1x1 relation – single row, single column
- Uses:
 - Compute an associated value

```
SELECT P.Name, (SELECT AVG(P1.Salary)
                  FROM Payroll AS P1
                  WHERE P.Job = P1.Job)
FROM Payroll AS P
```

Correlated subquery!
Inner query refers to
attributes from the
outer query

Subqueries in SELECT

- Must return a single value
 - A 1x1 relation – single row, single column
- Uses:
 - Compute an associated value

```
SELECT P.Name, (SELECT AVG(P1.Salary)
                  FROM Payroll AS P1
                  WHERE P.Job = P1.Job)
FROM Payroll AS P
```

Correlated subquery!
Semantics are that the
entire subquery is
recomputed for each
tuple

Subqueries in SELECT

For each person find the average salary of their job

```
SELECT P.Name, (SELECT AVG(P1.Salary)
                  FROM Payroll AS P1
                  WHERE P.Job = P1.Job)
FROM Payroll AS P
```



Same (decorrelated and unnested)

```
SELECT P1.Name, AVG(P2.Salary)
FROM Payroll AS P1, Payroll AS P2
WHERE P1.Job = P2.Job
GROUP BY P1.UserID, P1.Name
```

FWGHOS recall

```
SELECT P1.Name, AVG(P2.Salary)
FROM Payroll AS P1, Payroll AS P2
WHERE P1.Job = P2.Job
GROUP BY P1.UserID, P1.Name
```

P1

P2

UserID	Name	Job	Salary	UserID	Name	Job	Salary
123	Jack	TA	50000	123	Jack	TA	50000
123	Jack	TA	50000	345	Allison	TA	60000
345	Allison	TA	60000	345	Allison	TA	60000
345	Allison	TA	60000	123	Jack	TA	50000
567	Magda	Prof	90000	567	Magda	Prof	90000
567	Magda	Prof	90000	789	Dan	Prof	100000
789	Dan	Prof	100000	789	Dan	Prof	100000
789	Dan	Prof	100000	567	Magda	Prof	90000

FWGHOS recall

```
SELECT P1.Name, AVG(P2.Salary)
FROM Payroll AS P1, Payroll AS P2
WHERE P1.Job = P2.Job
GROUP BY P1.UserID, P1.Name
```

P1

P2

UserID	Name	Job	Salary	UserID	Name	Job	Salary
123	Jack	TA	50000	123	Jack	TA	50000
123	Jack	TA	50000	345	Allison	TA	60000
345	Allison	TA	60000	345	Allison	TA	60000
345	Allison	TA	60000	123	Jack	TA	50000
567	Magda	Prof	90000	567	Magda	Prof	90000
567	Magda	Prof	90000	789	Dan	Prof	100000
789	Dan	Prof	100000	789	Dan	Prof	100000
789	Dan	Prof	100000	567	Magda	Prof	90000

FWGHOS recall

```
SELECT P1.Name, AVG(P2.Salary)
```

```
FROM Payroll AS P1, Payroll
```

```
WHERE P1.Job = P2.Job
```

```
GROUP BY P1.UserID, P1.Name
```

p1.Name	AVG
Jack	55000
Allison	55000
Magda	95000
Dan	95000

P1

P2

UserID	Name	Job	Salary	UserID	Name	Job	Salary
123	Jack	TA	50000	123	Jack	TA	50000
123	Jack	TA	50000	345	Allison	TA	60000
345	Allison	TA	60000	345	Allison	TA	60000
345	Allison	TA	60000	123	Jack	TA	50000
567	Magda	Prof	90000	567	Magda	Prof	90000
567	Magda	Prof	90000	789	Dan	Prof	100000
789	Dan	Prof	100000	789	Dan	Prof	100000
789	Dan	Prof	100000	567	Magda	Prof	90000

Subqueries in SELECT

For each person find the number of cars they drive

```
SELECT P.Name, (SELECT COUNT(R.Car)
                  FROM Regist AS R
                  WHERE P.UserID =
                      R.UserID)
FROM Payroll AS P
```



Same? **Discuss!**

```
SELECT P.Name, COUNT(R.Car)
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID
GROUP BY P.UserID, P.Name
```

Subqueries in SELECT

For each person find the number of cars they drive

```
SELECT P.Name, (SELECT COUNT(R.Car)
                  FROM Regist AS R
                  WHERE P.UserID =
                      R.UserID)
FROM Payroll AS P
```

0-count case not covered!

```
SELECT P.Name, COUNT(R.Car)
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID
GROUP BY P.UserID, P.Name
```

Subqueries in SELECT

For each person find the number of cars they drive

```
SELECT P.Name, (SELECT COUNT(R.Car)
                  FROM Regist AS R
                  WHERE P.UserID =
                      R.UserID)
FROM Payroll AS P
```



Still possible to decorrelate and
unnest

Subqueries in SELECT

For each person find the number of cars they drive

```
SELECT P.Name, (SELECT COUNT(R.Car)
                  FROM Regist AS R
                  WHERE P.UserID =
                      R.UserID)
FROM Payroll AS P
```



Still possible to decorrelate and
unnest

```
SELECT P.Name, COUNT(R.Car)
FROM Payroll AS P LEFT OUTER JOIN
    Regist AS R ON P.UserID = R.UserID
GROUP BY P.UserID, P.Name
```

The Witnessing Problem

- Also known as argmax/argmin
- Ex: Return the person with the highest salary for each job type

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

The Witnessing Problem

- Also known as argmax/argmin
- Ex: Return the person with the highest salary for each job type

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

```
SELECT Name, MAX(Salary)
FROM Payroll
GROUP BY Job
```

Easy right?

The Witnessing Problem

- Also known as argmax/argmin
- Ex: Return the person with the highest salary for each job type

UserID		Salary
123		50000
345		60000
567		90000
789		100000

```
SELECT Name, MAX (Salary)  
FROM Payroll  
GROUP BY Job
```

The Witnessing Problem

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Name	MAX(Salary)
???	60000
???	100000



```
SELECT Name, MAX(Salary)
FROM Payroll
GROUP BY Job
```


The Witnessing Problem

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Return the person with the highest salary for each job type

How do we witness the maxima for a group?

Discuss!

Conceptual ideas are great

The Witnessing Problem

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Return the person with the highest salary for each job type

Main idea:

we need to join the respective maxima to each row

The Witnessing Problem

UserID	Name	Job	Salary	maxima
123	Jack	TA	50000	60000
345	Allison	TA	60000	60000
567	Magda	Prof	90000	100000
789	Dan	Prof	100000	100000

Return the person with the highest salary for each job type

Main idea:

we need to join the respective maxima to each row

Subqueries in FROM

- Uses:

- Solve subproblems that can be later joined/evaluated

```
SELECT P.Name, P.Salary
FROM Payroll AS P,
      (SELECT P1.Job AS Job,
          MAX(P1.Salary) AS Salary
       FROM Payroll AS P1
       GROUP BY P1.Job) AS Pmax
WHERE P.Job = Pmax.Job AND
      P.Salary = Pmax.Salary
```

Subqueries in FROM

- Equivalent to a WITH subquery

```
WITH MaxPay AS
    (SELECT P1.Job AS Job,
           MAX(P1.Salary) AS Salary
     FROM Payroll AS P1
     GROUP BY P1.Job)
SELECT P.Name, P.Salary
FROM Payroll AS P, MaxPay AS MP
WHERE P.Job = MP.Job AND
       P.Salary = MP.Salary
```



Subqueries in WHERE/HAVING

- Can return a single value
- Uses:
 - Compare with another value

```
SELECT P.Name, P.Salary
FROM Payroll AS P
WHERE P.Salary =
    (SELECT MAX(P1.Salary) AS Salary
     FROM Payroll AS P1
     WHERE P1.Job = P.Job)
```

Subqueries in WHERE/HAVING

- Can return a single value
- Uses:
 - Compare with another value

```
SELECT P.Name, P.Salary
FROM Payroll AS P
WHERE P.Salary =
    (SELECT MAX(P1.Salary) AS Salary
     FROM Payroll AS P1
     WHERE P1.Job = P.Job)
```

**Correlated subquery
alert!**

Witnessing Unnested

UserID	Name	Job	Salary	maxima
123	Jack	TA	50000	60000
345	Allison	TA	60000	60000
567	Magda	Prof	90000	100000
789	Dan	Prof	100000	100000

Return the person with the highest salary for each job type

```
SELECT P.Name,  
        MAX(Pmax.Salary) as MaxSalary  
FROM Payroll AS P, Payroll AS Pmax  
WHERE P.Job = Pmax.Job  
GROUP BY Pmax.Job, P.Name, P.Salary  
HAVING P.Salary = MaxSalary
```


Witnessing Unnested

```
SELECT P1.Name, MAX(Pmax.Salary)
FROM Payroll AS P1, Payroll AS Pmax
WHERE P1.Job = Pmax.Job
GROUP BY Pmax.Job, P1.Salary, P1.Name
HAVING P1.Salary = MAX(Pmax.Salary)
```

P1				Pmax			
UserID	Name	Job	Salary	UserID	Name	Job	Salary
123	Jack	TA	50000	123	Jack	TA	50000
123	Jack	TA	50000	345	Allison	TA	60000
345	Allison	TA	60000	345	Allison	TA	60000
345	Allison	TA	60000	123	Jack	TA	50000
567	Magda	Prof	90000	567	Magda	Prof	90000
567	Magda	Prof	90000	789	Dan	Prof	100000
789	Dan	Prof	100000	789	Dan	Prof	100000
789	Dan	Prof	100000	567	Magda	Prof	90000

Subqueries in WHERE/HAVING

- Can return a relation
- Uses:
 - Use with an existential or universal quantifier
 - (NOT) EXISTS, (NOT) IN, ANY, ALL

```
SELECT ..... WHERE EXISTS (subquery);  
SELECT ..... WHERE NOT EXISTS (subquery);  
SELECT ..... WHERE attr IN (subquery);  
SELECT ..... WHERE attr NOT IN (subquery);  
SELECT ..... WHERE const > ANY (subquery);  
SELECT ..... WHERE const > ALL (subquery);
```

Subqueries in WHERE/HAVING

- Can return a relation
- Uses:
 - Use with an existential or universal quantifier
 - (NOT) EXISTS, (NOT) IN, ANY, ALL

Ex: Find all people who drive some car made before 2017.



Existential quantifier

Subqueries in WHERE/HAVING

- Can return a relation
- Uses:
 - Use with an existential or universal quantifier
 - (NOT) EXISTS, (NOT) IN, ANY, ALL

Ex: Find all people who drive some car made before 2017.

EXISTS (subquery) returns true iff
cardinality of subquery > 0

```
SELECT P.Name
FROM Payroll AS P
WHERE EXISTS (SELECT *
                FROM Regist R
                WHERE R.UserID = P.UserID
                       AND R.Year < 2017)
```

Subqueries in WHERE/HAVING

- Can return a relation
- Uses:
 - Use with an existential or universal quantifier
 - (NOT) EXISTS, (NOT) IN, ANY, ALL

Ex: Find all people who drive some car made before 2017.

attr IN (subquery) returns true iff
value of attr is contained in subquery

```
SELECT P.Name
FROM Payroll AS P
WHERE P.UserID IN (SELECT R.UserID
                     FROM Regist R
                     WHERE R.Year < 2017)
```

Subqueries in WHERE/HAVING

- Can return a relation
- Uses:
 - Use with an existential or universal quantifier
 - (NOT) EXISTS, (NOT) IN, ANY, ALL

Ex: Find all people who drive some car made before 2017.

attr IN (subquery) returns true iff
value of attr is in the result of subquery

Decorrelated!

```
SELECT P.Name
FROM Payroll AS P
WHERE P.UserID IN (SELECT R.UserID
                     FROM Regist R
                     WHERE R.Year < 2017)
```

Subqueries in WHERE/HAVING

- Can return a relation
- Uses:
 - Use with an existential or universal quantifier
 - (NOT) EXISTS, (NOT) IN, ANY, ALL

Ex: Find all people who drive some car made before 2017.

const > ANY (sub) returns true iff
const > value for at least one value in sub


```
SELECT P.Name
FROM Payroll AS P
WHERE 2017 > ANY (SELECT R.Year
                   FROM Regist R
                   WHERE
                       P.UserID = R.UserID)
```

Subqueries in WHERE/HAVING

- Can return a relation
- Uses:
 - Use with an existential or universal quantifier
 - (NOT) EXISTS, (NOT) IN, ANY, ALL

Ex: Find all people who drive some car made before 2017.

```
SELECT P.Name
FROM Payroll A
WHERE 2017 > ANY (SELECT R.Year
                    FROM Regist R
                    WHERE
                        P.UserID = R.UserID)
```



Subqueries in WHERE/HAVING

- Can return a relation
- Uses:
 - Use with an existential or universal quantifier
 - (NOT) EXISTS, (NOT) IN, ANY, ALL

Ex: Find all people who drive some car made before 2017.

Unnesting existential quantifiers is easy!

```
SELECT DISTINCT P.Name
FROM Payroll AS P, Regist R
WHERE P.UserID = R.UserID AND
      R.Year < 2017
```



Subqueries in WHERE/HAVING

- Can return a relation
- Uses:
 - Use with an existential or universal quantifier
 - (NOT) EXISTS, (NOT) IN, ANY, ALL

Ex: Find all people who drive some car made before 2017.

```
SELECT DISTINCT P.Name
FROM Payroll AS P, Regist R
WHERE P.UserID = R.UserID AND
      R.Year < 2017
```

Unnesting existential
quantifiers is easy!



Expert SQL style

Subqueries in WHERE/HAVING

- Can return a relation
- Uses:
 - Use with an existential or universal quantifier
 - (NOT) EXISTS, (NOT) IN, ANY, ALL

Ex: Find all people who drive only cars older than 2017.



Universal quantifier

Subqueries in WHERE/HAVING

- Can return a relation
- Uses:
 - Use with an existential or universal quantifier
 - (NOT) EXISTS, (NOT) IN, ANY, ALL

Ex: Find all people who drive only cars older than 2017.

Universal quantifier

Not easy :(

Subqueries in WHERE/HAVING

- Can return a relation
- Uses:
 - Use with an existential or universal quantifier
 - (NOT) EXISTS, (NOT) IN, ANY, ALL

Ex: Find all people who drive older cars than 2017.

Find all the other people, the ones who DO drive newer cars

```
(SELECT R.UserID
FROM Regist AS R
WHERE R.Year >= 2017)
```

Subqueries in WHERE/HAVING

- Can return a relation
- Uses:
 - Use with an existential or universal quantifier
 - (NOT) EXISTS, (NOT) IN, ANY, ALL

Ex: Find all people who drive older cars than 2017.

```
SELECT P.Name
FROM Payroll AS P
WHERE P.UserID NOT IN (SELECT R.UserID
                        FROM Regist AS R
                        WHERE R.Year >= 2017)
```

Find all the other people, the ones who DO drive newer cars

Subqueries in WHERE/HAVING

- Can return a relation
- Uses:
 - Use with an existential or universal quantifier
 - (NOT) EXISTS, (NOT) IN, ANY, ALL

Ex: Find all people who drive only cars older than 2017.

```
SELECT P.Name
FROM Payroll AS P
WHERE NOT EXISTS (SELECT *
                    FROM Regist AS R
                    WHERE R.Year >= 2017
                    AND P.UserID = R.UserID)
```

Subqueries in WHERE/HAVING

- Can return a relation
- Uses:
 - Use with an existential or universal quantifier
 - (NOT) EXISTS, (NOT) IN, ANY, ALL

Ex: Find all people who drive only cars older than 2017.

const > ALL(sub) returns true iff
const > value for all values in sub

```
SELECT P.Name
FROM Payroll AS P
WHERE 2017 > ALL (SELECT R.Year
                    FROM Regist AS R
                    WHERE
                        P.UserID = R.UserID)
```


Subqueries in WHERE/HAVING

- Can return a relation
- Uses:
 - Use with an existential or universal quantifier
 - (NOT) EXISTS, (NOT) IN, ANY, ALL

Ex: Find all people who drive only cars older than 2017.

```
SELECT P.Name
FROM Payroll A
WHERE 2017 > ALL (SELECT R.Year
                    FROM Regist AS R
                    WHERE
                        P.UserID = R.UserID)
```

Not supported in sqlite :(

Unnesting

Can we unnest the universal quantifier query?

First, a discussion on the concept of monotonicity....


Monotonicity

Monotone

A **Monotonic** query is one that obeys the following rule where I and J are data instances and q is a query:

$$I \subseteq J \rightarrow q(I) \subseteq q(J)$$

That is for any superset of I , the query over that superset must contain at least the query results of I .




Monotonicity

Monotone

A **Monotonic** query is one that obeys the following rule where I and J are data instances and q is a query:

$$I \subseteq J \rightarrow q(I) \subseteq q(J)$$

That is for any superset of I , the query over that superset must contain at least the query results of I .



Monotone queries can be similar to monotonically increasing functions when considering cardinalities of results

Monotonicity

Monotone

A **Monotonic** query is one that obeys the following rule where I and J are data instances and q is a query:

$$I \subseteq J \rightarrow q(I) \subseteq q(J)$$

That is for any superset of I, the query over that superset must contain at least the query results of I.

```
SELECT P.Name, P.Car  
      FROM Payroll AS P, Regist AS R  
      WHERE P.UserID = R.UserID
```

Is this query monotone?

Monotonicity

Monotone

A **Monotonic** query is one that obeys the following rule where I and J are data instances and q is a query:

$$I \subseteq J \rightarrow q(I) \subseteq q(J)$$

That is for any superset of I, the query over that superset must contain at least the query results of I.

```
SELECT P.Name, P.Car  
      FROM Payroll AS P, Regist AS R  
      WHERE P.UserID = R.UserID
```

Is this query monotone? **Yes!**

Monotonicity

Monotone

A **Monotonic** query is one that obeys the following rule where I and J are data instances and q is a query:

$$I \subseteq J \rightarrow q(I) \subseteq q(J)$$

That is for any superset of I, the query over that superset must contain at least the query results of I.

I can't add tuples to Payroll or Regist that would "remove" a previous result

```
SELECT P.Name, P.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID
```

Is this query monotone? **Yes!**

Monotonicity

Monotone

A **Monotonic** query is one that obeys the following rule where I and J are data instances and q is a query:

$$I \subseteq J \rightarrow q(I) \subseteq q(J)$$

That is for any superset of I, the query over that superset must contain at least the query results of I.

```
SELECT P.Name  
  FROM Payroll AS P  
 WHERE P.Salary >= ALL (SELECT Salary  
                        FROM Payroll)
```

Is this query monotone?

Monotonicity

Monotone

A **Monotonic** query is one that obeys the following rule where I and J are data instances and q is a query:

$$I \subseteq J \rightarrow q(I) \subseteq q(J)$$

That is for any superset of I, the query over that superset must contain at least the query results of I.

```
SELECT P.Name  
  FROM Payroll AS P  
 WHERE P.Salary >= ALL (SELECT Salary  
                        FROM Payroll)
```

Is this query monotone? **No!**

Monotonicity

Monotone

A **Monotonic** query is one that obeys the following rule where I and J are data instances and q is a query:

$$I \subseteq J \rightarrow q(I) \subseteq q(J)$$

That is for any superset of I, the query over that superset must contain at least the query results of I.

```
SELECT P.Name  
  FROM Payroll AS P  
 WHERE P.Salary >= ALL (SELECT Salary  
                        FROM Payroll)
```

I can add a tuple to Payroll that has a higher salary value than any other

Is this query monotone? **No!**

Monotonicity

Monotone

A **Monotonic** query is one that obeys the following rule where I and J are data instances and q is a query:

$$I \subseteq J \rightarrow q(I) \subseteq q(J)$$

That is for any superset of I, the query over that superset must contain at least the query results of I.

```
SELECT P.Job, COUNT (*)  
  FROM Payroll AS P  
GROUP BY P.Job
```

Is this query monotone?

Monotonicity

Monotone

A **Monotonic** query is one that obeys the following rule where I and J are data instances and q is a query:

$$I \subseteq J \rightarrow q(I) \subseteq q(J)$$

That is for any superset of I, the query over that superset must contain at least the query results of I.

```
SELECT P.Job, COUNT (*)  
  FROM Payroll AS P  
GROUP BY P.Job
```

Is this query monotone? **No!**

Monotonicity

Monotone

A **Monotonic** query is one that obeys the following rule where I and J are data instances and q is a query:

$$I \subseteq J \rightarrow q(I) \subseteq q(J)$$

That is for any superset of I, the query over the superset contains at least the query results of I.

Aggregates generally are sensitive to any new tuples since the aggregate value will change

```
SELECT P.Job, COUNT (*)  
FROM Payroll AS P  
GROUP BY P.Job
```

Is this query monotone? **No!**

Monotonicity

Theorem:

If Q is a SELECT-FROM-WHERE query that does not have subqueries or aggregates, then it is monotone.

Monotonicity

Theorem:

If Q is a SELECT-FROM-WHERE query that does not have subqueries or aggregates, then it is monotone.

Proof:

We use nested loop semantics. If we insert a tuple in relation R, this will not remove any tuples from the answer.

```
SELECT a1, a2, ..., ak  
FROM R1 AS x1, R2 AS x2, ..., Rn AS xn  
WHERE Conditions
```

```
for x1 in R1 do  
  for x2 in R2 do  
    ...  
    for xn in Rn do  
      if Conditions  
        output (a1,...,ak)
```

Monotonicity

Theorem:

The query “Find all people who drive only cars older than 2017” is not monotone.

Proof:

We use example. For user 123 who previously only drove a car made in 2009, we add another car made in 2018. Now user 123 does not appear in the results. Thus, the query is not monotone.

Monotonicity

Theorem:

The query “Find all people who drive only cars older than 2017” is not monotonic.

Proof:

We use a query that is not previously seen. We add another car to the database. Now user 123 does not own a car. Thus, the query is not monotonic.

If a query is not monotonic, then we can't write it as a SELECT-FROM-WHERE query without subqueries

Queries That Cannot Be S-F-W

- Queries with universal quantifiers or negation

```
SELECT P.Name
FROM Payroll AS P
WHERE P.UserID NOT IN (SELECT R.UserID
                        FROM Regist AS R
                        WHERE R.Year < 2017)
```

```
SELECT P.Name
FROM Payroll AS P
WHERE P.Salary >= ALL (SELECT Salary
                        FROM Payroll)
```


Bonus: Set Operations

- SQL mimics set theory in many ways
 - Bag = duplicates allowed
 - **UNION (ALL)** □ set union (bag union)
 - **INTERSECT (ALL)** □ set intersection (bag intersection)
 - **EXCEPT (ALL)** □ set difference (bag difference)
- SQL Server Management Studio 2017
 - INTERSECT ALL not supported
 - EXCEPT ALL not supported



Set Operations

- SQL set-like operators basically slap two queries together (not really a subquery...)



Takeaways

- Subqueries let us express some problems more easily
- Many subqueries can be unnested
- Some cannot be (think non-monotonic queries – universal quantifiers, negation, aggregates)
- SQL set operations behave much like set theory