

# Introduction to Data Management

## Transactions: Serializability

Alyssa Pittman

Based on slides by Jonathan Leang, Dan Suciu, et al

Paul G. Allen School of Computer Science and Engineering  
University of Washington, Seattle

# Outline

- Concurrency control problems
- Transactions
- Serializable schedules
- Conflict serializability

# Transactions

How do we support multiple people using a database at the same time?

- Multiple application users
- Multiple application programmers
- Multiple analysts


Imagine a world where each person had to wait in line to use your database 😞

# Concurrency Control Problems

- Non-Atomic Operations
- Lost Update
- Dirty/Inconsistent Read
- Unrepeatable Read
- Phantom Read

# Non-Atomic Operations


- > I'll book this airline flight!
- > ...submit credit card number
- > loading.gif
- > ...bank account goes down...
- > no booking confirmation
- > waiting.png
- > browser crashes
- > still no confirmation



# Non-Atomic Operations

- > I'll book this airline flight!
- > ...submit credit card number
- > loading.gif
- > ...bank account goes down...
- > no booking confirmation
- > waiting.png
- > browser crashes
- > still no confirmation

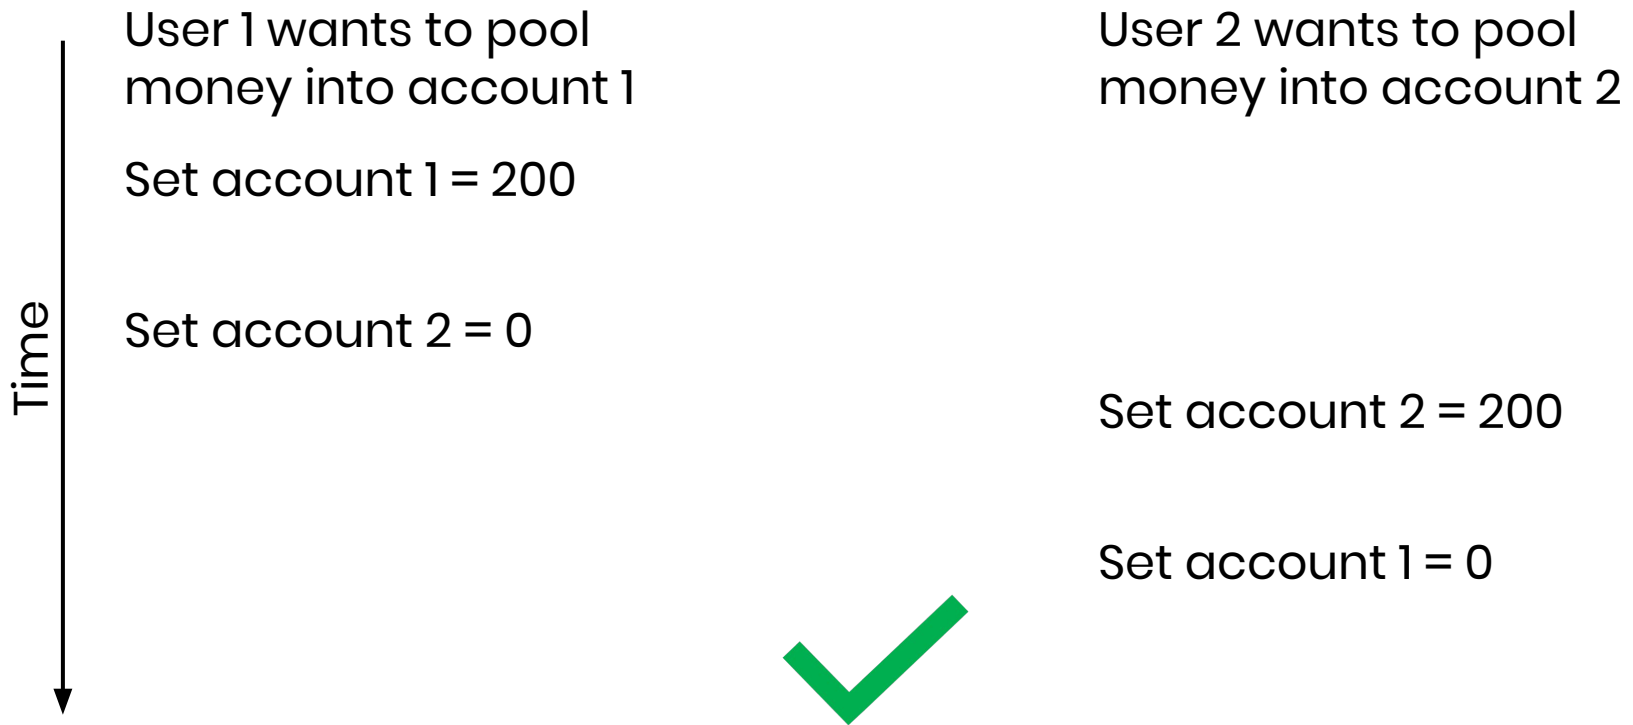
I would have been  
fine if nothing  
happened or if  
everything worked!



# Lost Update

- Write-Write (WW) conflict
- Consolidation scenario:

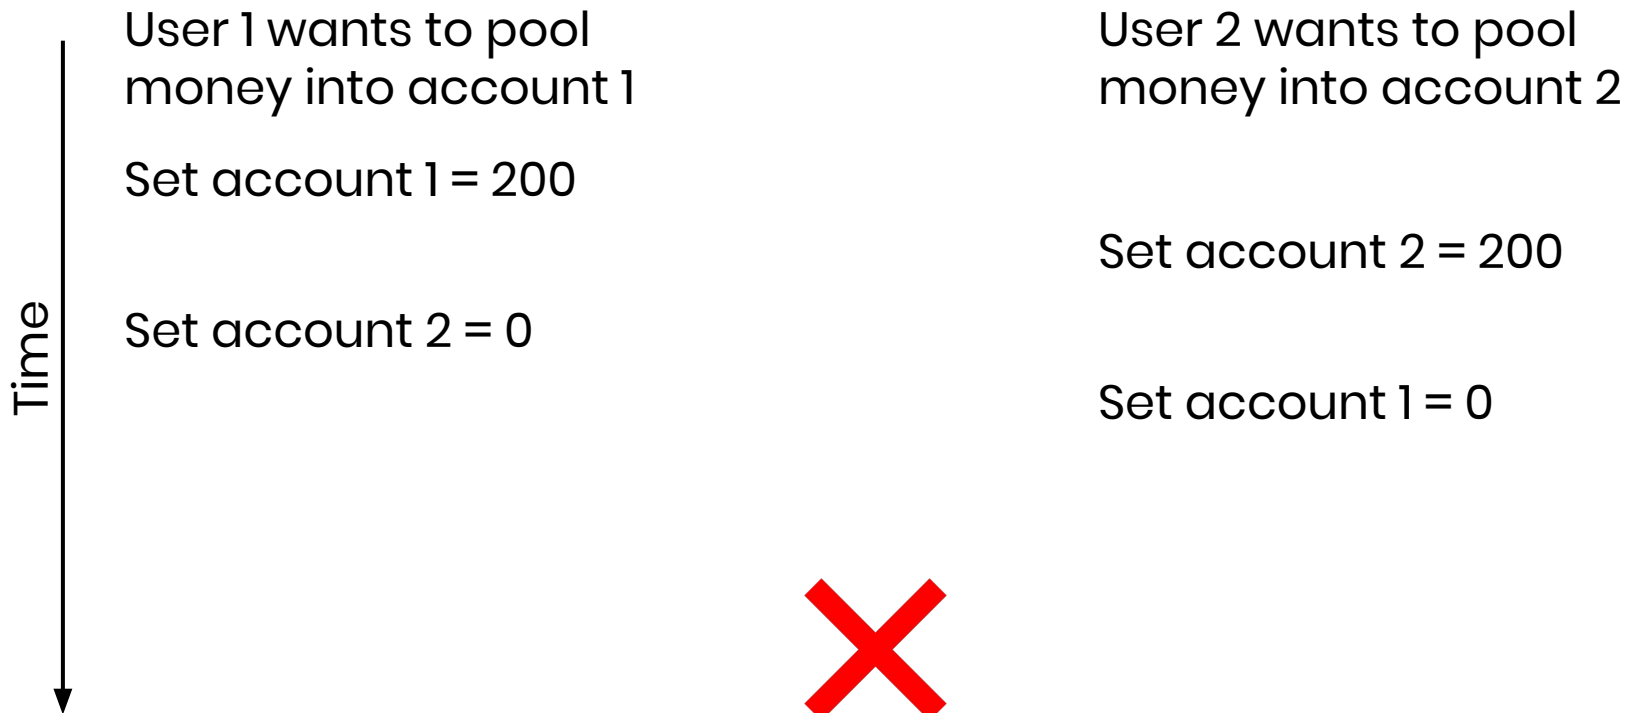
Account 1 = 100, Account 2 = 100



# Lost Update

- Write-Write (WW) conflict
- Consolidation scenario:

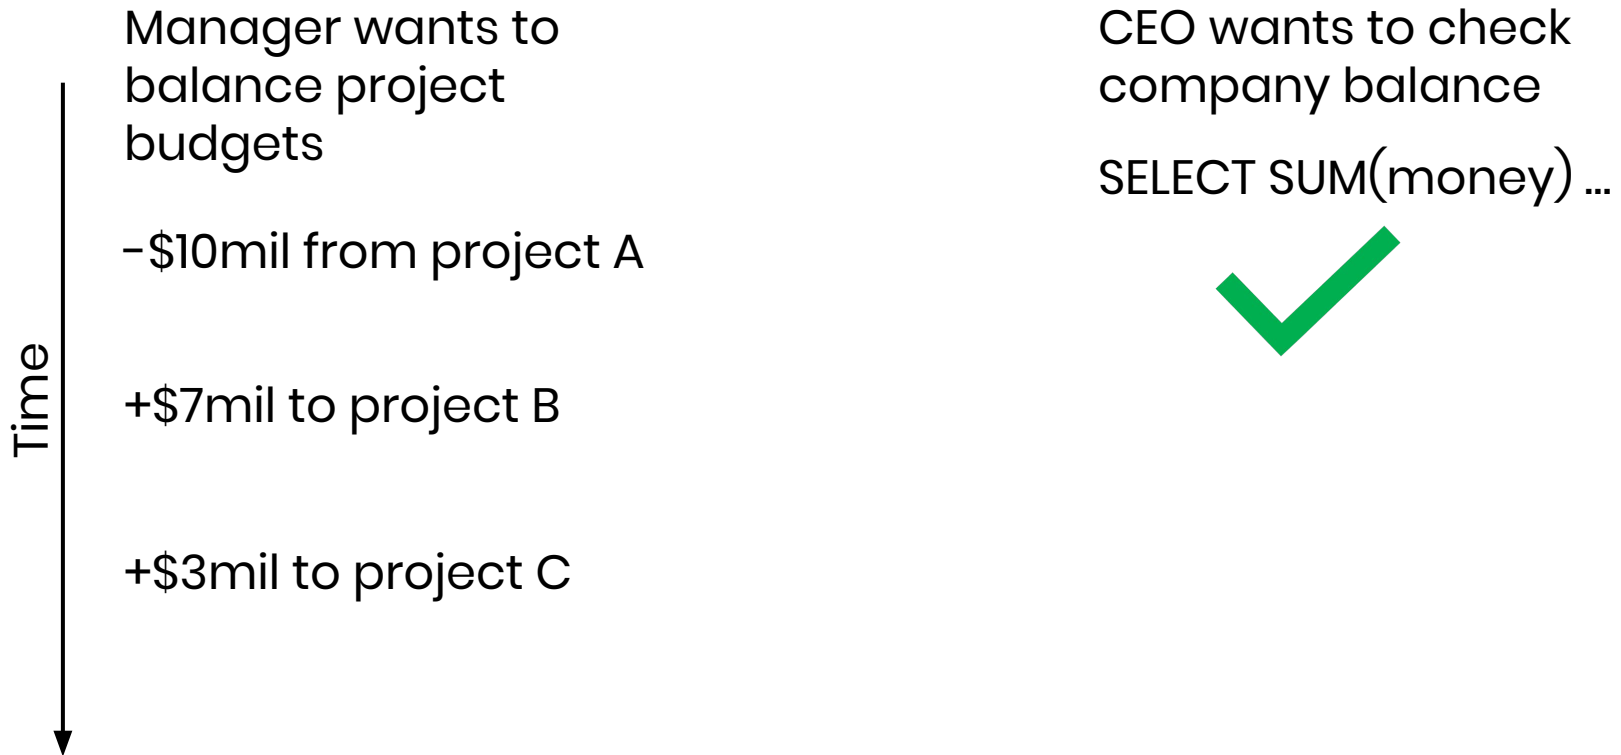
Account 1 = 100, Account 2 = 100





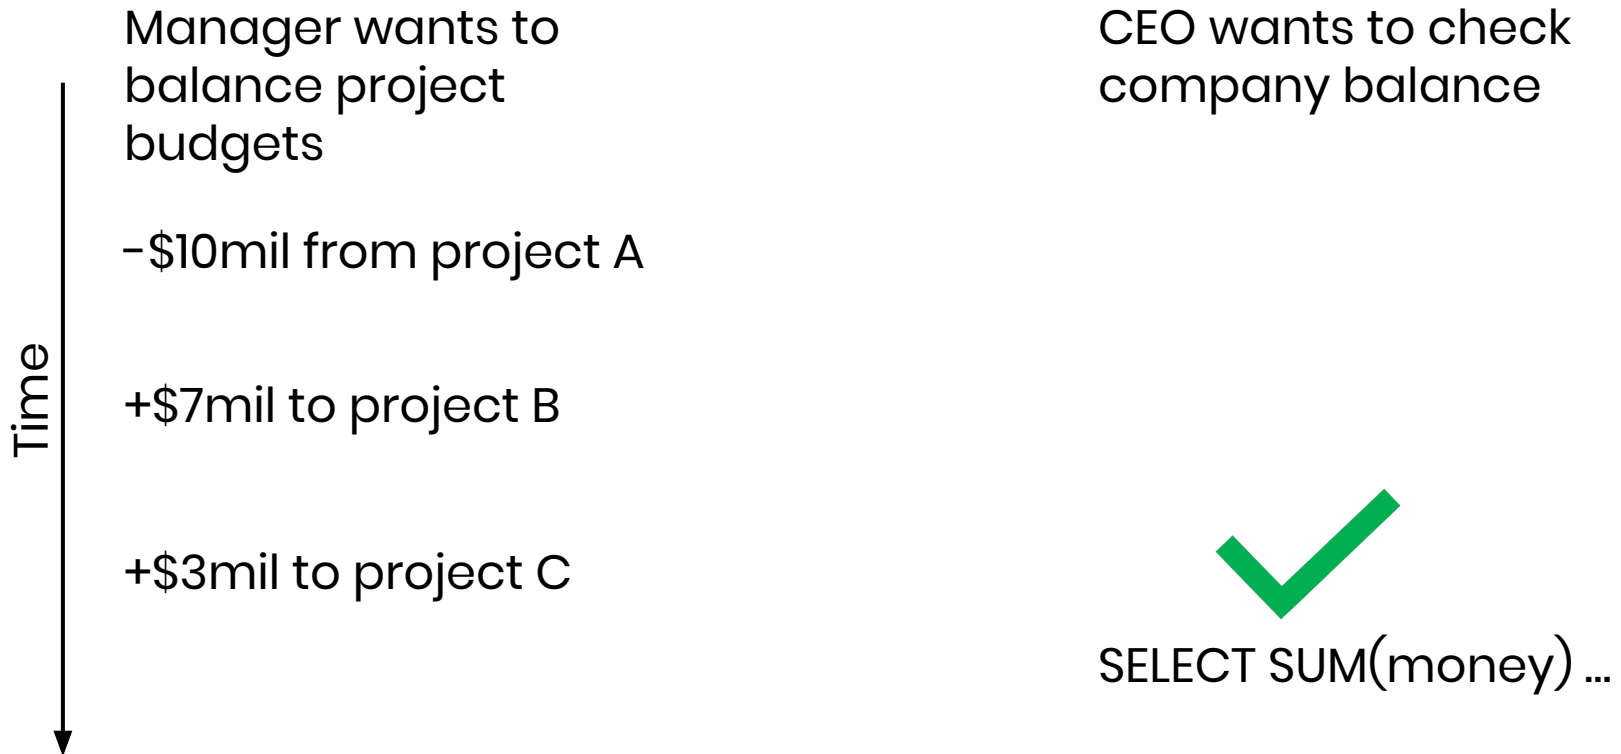
# Dirty/Inconsistent Read

- Write-Read (WR) conflict
- Budget management scenario:



# Dirty/Inconsistent Read

- Write-Read (WR) conflict
- Budget management scenario:



# Dirty/Inconsistent Read

- Write-Read (WR) conflict
- Budget management scenario:

Time ↓

Manager wants to  
balance project  
budgets


-\$10mil from project A

+\$7mil to project B

+\$3mil to project C

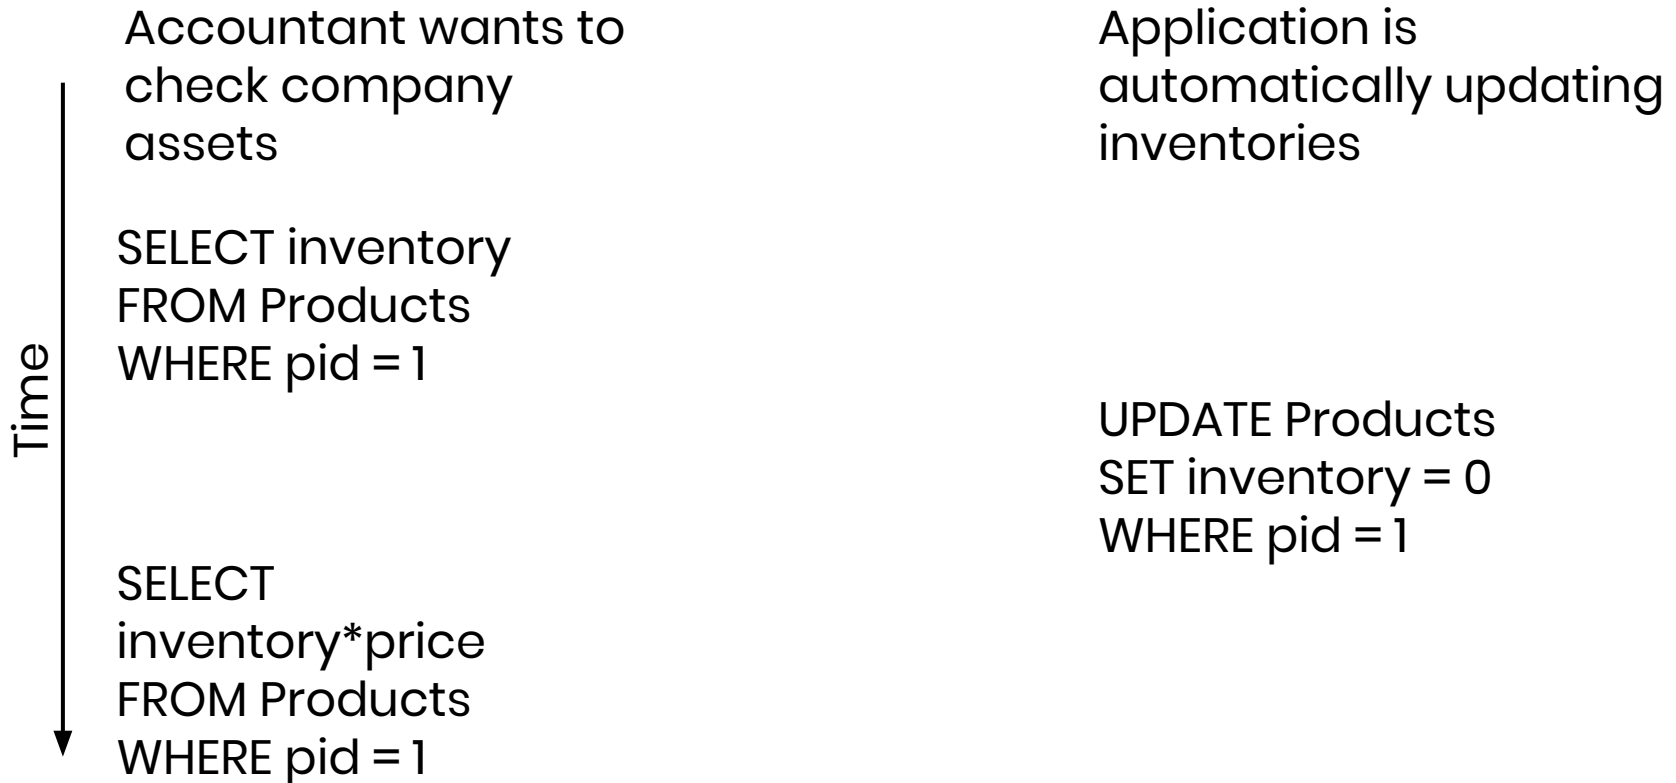
CEO wants to check  
company balance

SELECT SUM(money) ...



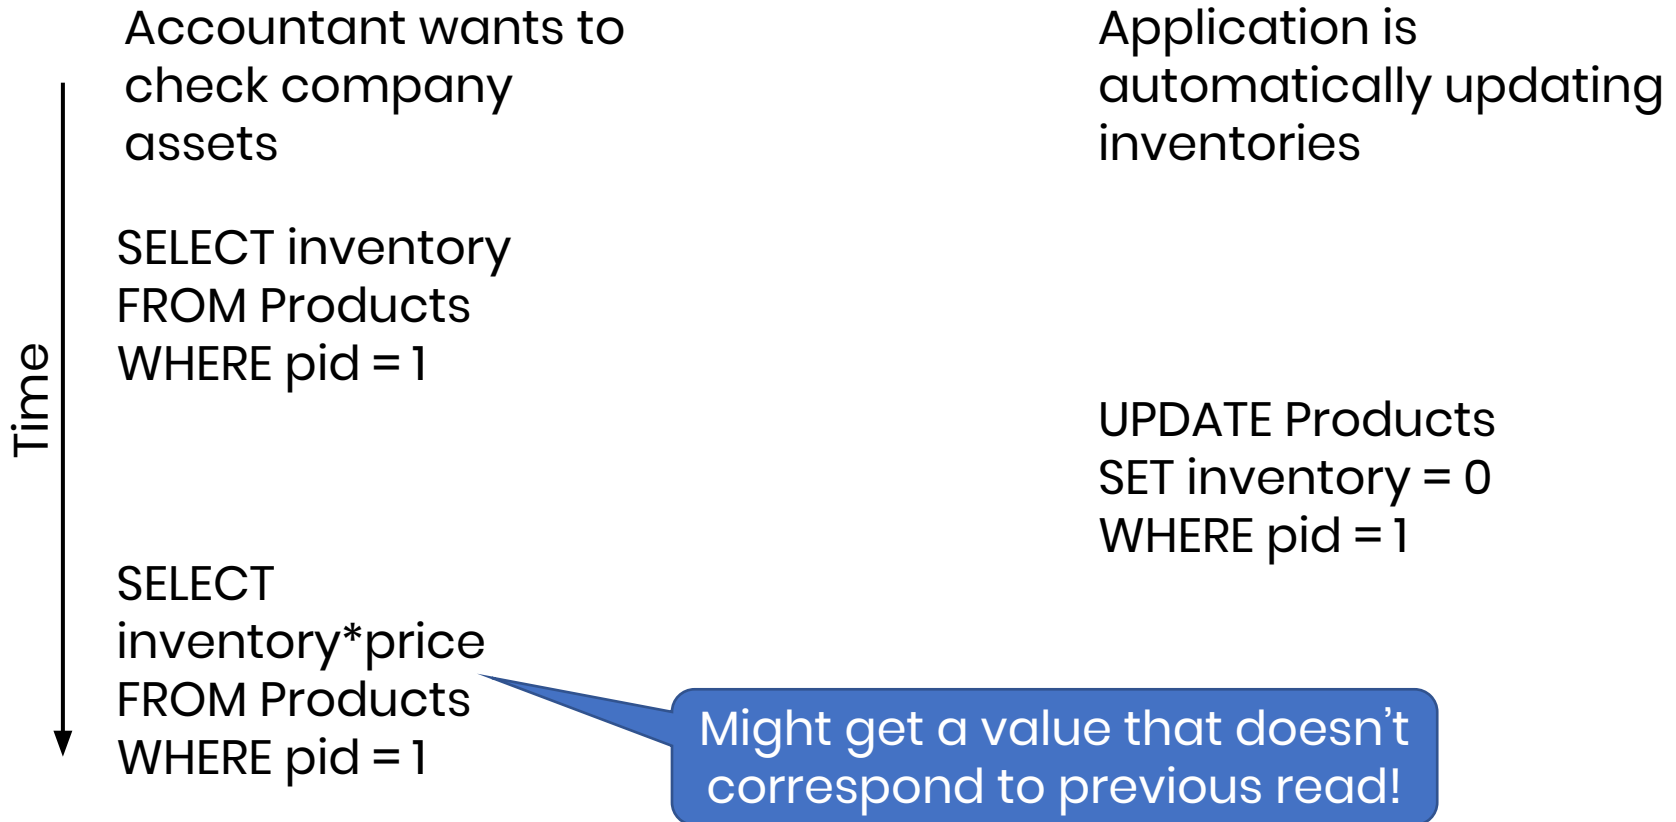
# Unrepeatable Read

- Read-Write (RW) conflict
- Asset checking scenario:



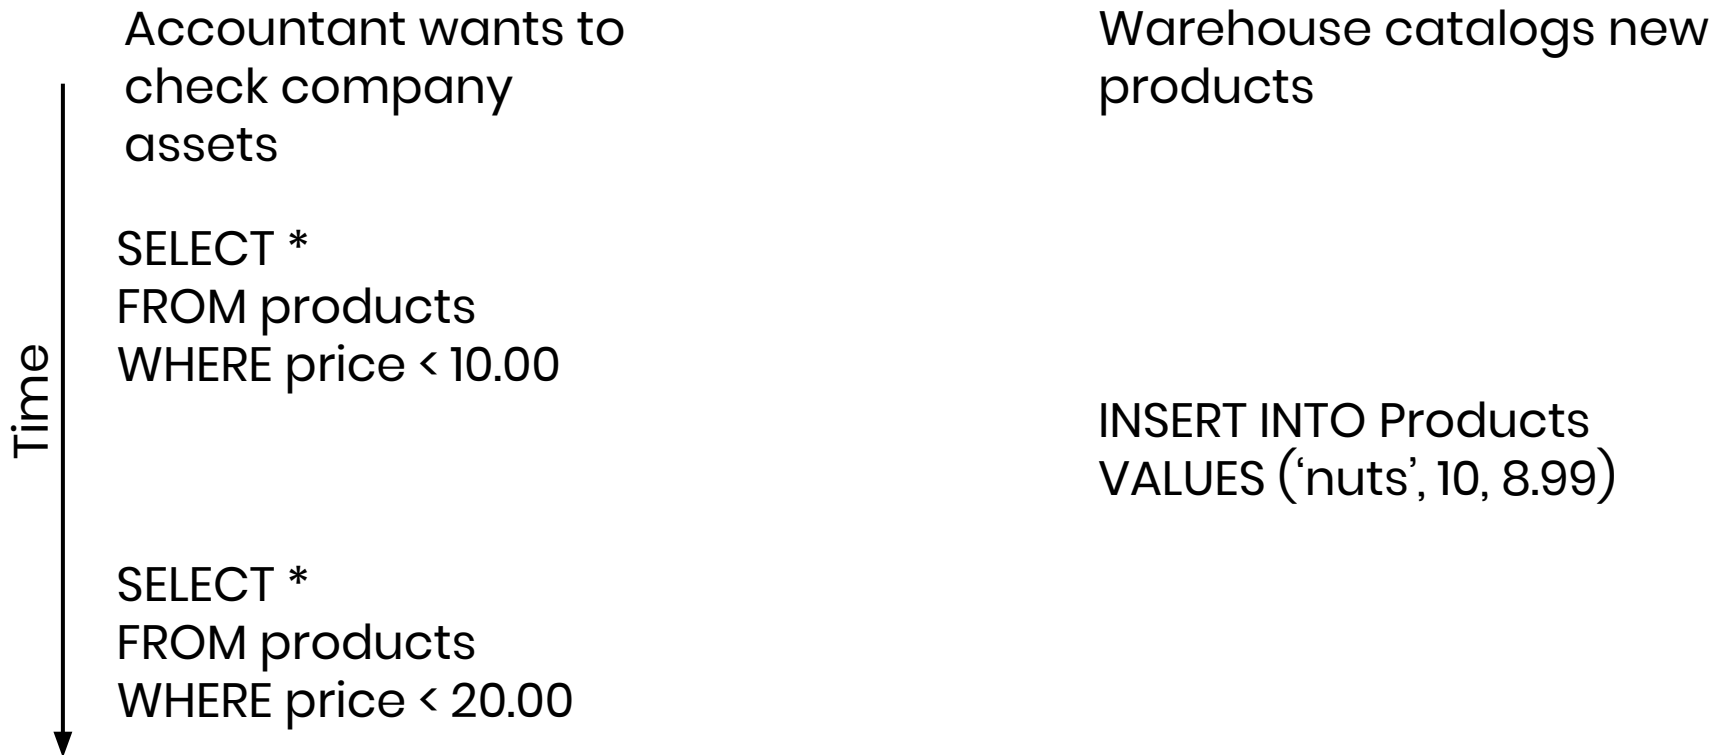
# Unrepeatable Read

- Read-Write (RW) conflict
- Asset checking scenario:



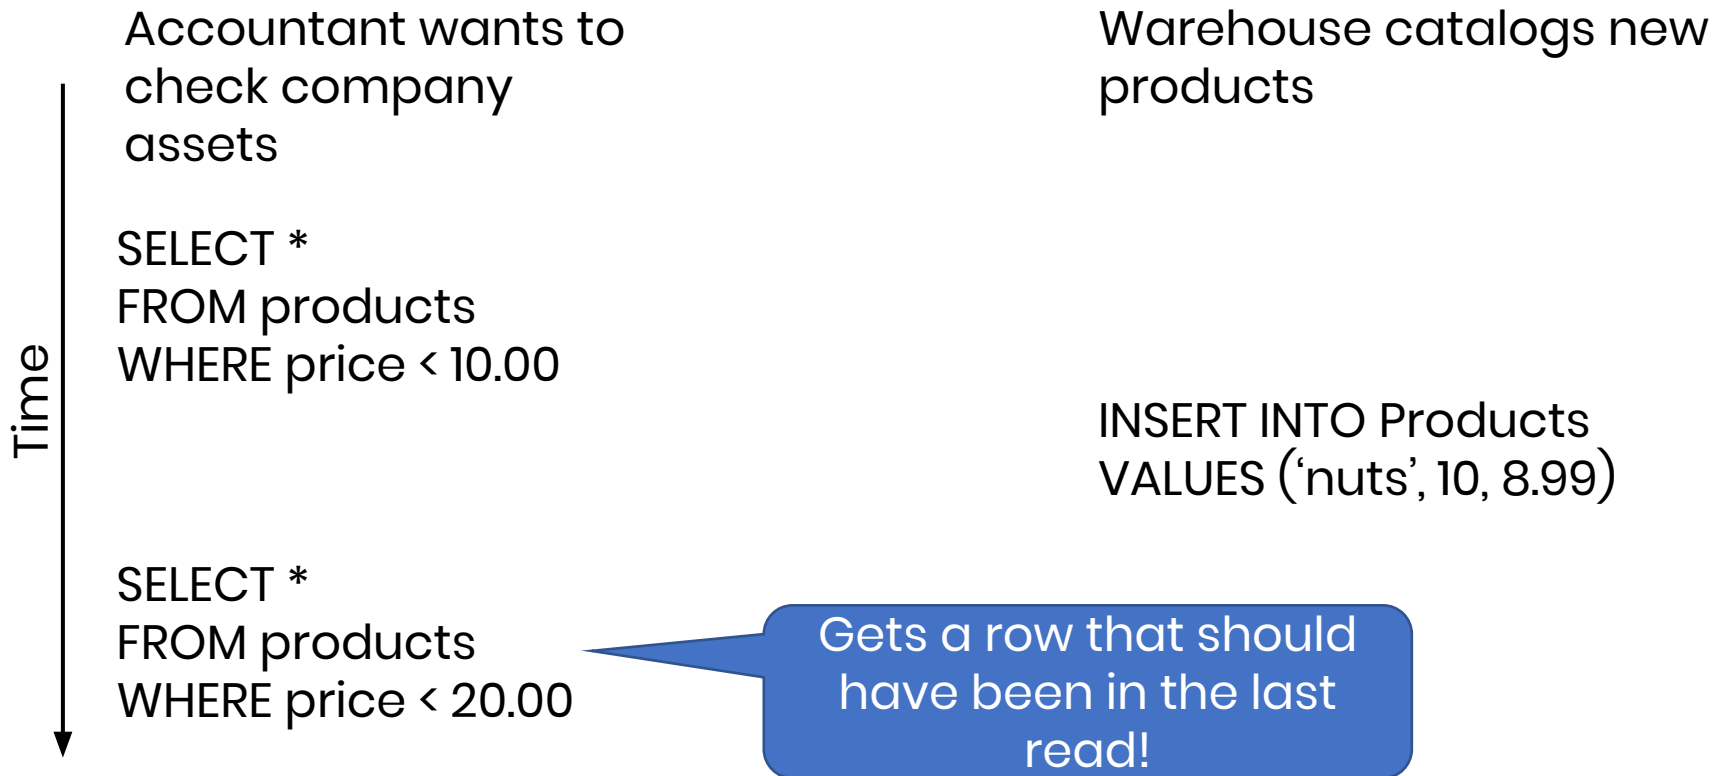
# Phantom Read

- Same read has more rows
- Asset checking scenario:



# Phantom Read

- Same read has more rows
- Asset checking scenario:



# ACID

- Atomic
  - Consistent
  - Isolated
  - Durable
- 
- Ideally a DBMS follows these principles, but sacrificing some behavior for performance gains is common
  - Definitely needs to follow these principles if you are dealing with \$\$\$



# Atomic

- Operation encapsulation
- A transaction is atomic if the state shows all the effects of a transaction, or none of them (all or nothing)

# Consistent

- Integrity constraints and application specification
- A transaction begins with a valid database state and ends with a valid database state

# Isolated

- Concurrency management
- Running transactions has the same effect as if we ran each transaction one after another

# Durable

- Crash recovery
- Once a transaction has been committed, its effects remain in the database
- CSE 444 topic, not discussed in this class

# Transactions

- A collection of statements that are executed atomically (logically speaking!)
  - A single application function may involve multiple different operations
  - Transactions let them execute properly **together as if it were a single action**

```
BEGIN TRANSACTION  
[SQL Statements]  
COMMIT -- finalizes execution
```

```
BEGIN TRANSACTION  
[SQL Statements]  
ROLLBACK -- undo everything
```

```
[Single SQL Statement]
```

# Transactions

- A collection of statements that are executed atomically (logically speaking!)
  - A single application function may involve multiple different operations
  - Transactions let them execute properly **together as if it were a single action**

```
BEGIN TRANSACTION
[SQL Statements]
COMMIT -- finalizes execution
```

```
BEGIN TRANSACTION
[SQL Statements]
ROLLBACK -- undo everything
```

[Single SQL Statement]

By default, without BEGIN, each statement is its own transaction

# Do I need to implement this?

- **DBMS concurrency control is all based on specification**
- Merely specifying what your transactions are is good enough for the DBMS to take care of it as a single unit

# Transaction Modeling

- Logical perspective □ a database is a set of sets/bags of tuples
- Design perspective □ a database is a schema that models information
- Physical perspective □ a database is a catalog of organized files
- Transaction perspective □ a database is a **collection of elements** that can be **written to** or **read from**
  - Element granularity can vary depending on DBMS and/or user specification
  - Transactions are sequences of element reads and/or writes



# Schedules

- Transactions are sequence of element reads and/or writes
  - $R_i(A)$   $\square$  transaction  $i$  **reads** element  $A$
  - $W_i(A)$   $\square$  transaction  $i$  **updates** element  $A$
  - $I_i(A)$   $\square$  transaction  $i$  **inserts** an element  $A$
  - $D_i(A)$   $\square$  transaction  $i$  **deletes** an element  $A$
- Schedules are a sequence of interleaved actions from all transactions

# Schedules

- Transactions are sequence of element reads and/or writes
  - $R_i(A)$   $\square$  transaction  $i$  **reads** element  $A$
  - $W_i(A)$   $\square$  transaction  $i$  **updates** element  $A$
  - $I_i(A)$   $\square$  transaction  $i$  **inserts** an element  $A$
  - $D_i(A)$   $\square$  transaction  $i$  **deletes** an element  $A$
- Schedules are a sequence of interleaved actions from all transactions

# Serial Schedules

- A **serial schedule** is a schedule where each transaction would be executed in some order


# Transaction Schedule

T1	T2
R(A)	R(A)
W(A)	W(A)
R(B)	R(B)
W(B)	W(B)

# Serial Schedule Example

- T1 then T2

$R_1(A), W_1(A), R_1(B), W_1(B), R_2(A), W_2(A), R_2(B), W_2(B)$



T1	T2
R(A)	
W(A)	
R(B)	
W(B)	
	R(A)
	W(A)
	R(B)
	W(B)

# Serial Schedule Example

- T2 then T1

$R_2(A), W_2(A), R_2(B), W_2(B), R_1(A), W_1(A), R_1(B), W_1(B)$

T1	T2
	R(A)
	W(A)
	R(B)
	W(B)
R(A)	
W(A)	
R(B)	
W(B)	

# Serial Schedules

- A **serial schedule** is a schedule where each transaction would be executed in some order
- A **serializable schedule** is a schedule that is equivalent to a serial schedule
  - If the schedule were executed and you were given a before and after, you would not be able to tell if the transactions were interleaved

# Serializable Schedules

- Serializable to T1 then T2

$R_1(A), W_1(A), R_2(A), W_2(A), R_1(B), W_1(B), R_2(B), W_2(B)$

T1	T2
R(A)	
W(A)	
	R(A)
	W(A)
R(B)	
W(B)	
	R(B)
	W(B)



# Serializable Schedules

- Serializable to T1 then T2

$R_1(A), W_1(A), R_2(A), W_2(A), R_1(B), W_1(B), R_2(B), W_2(B)$

T1	T2
R(A)	
W(A)	
	R(A)
	W(A)
R(B)	
W(B)	
	R(B)
	W(B)

Looks like T2  
finished after T1  
for each element

# Serializable Schedules

- Not serializable to either order

$R_1(A), W_1(A), R_2(A), W_2(A), R_2(B), W_2(B), R_1(B), W_1(B)$

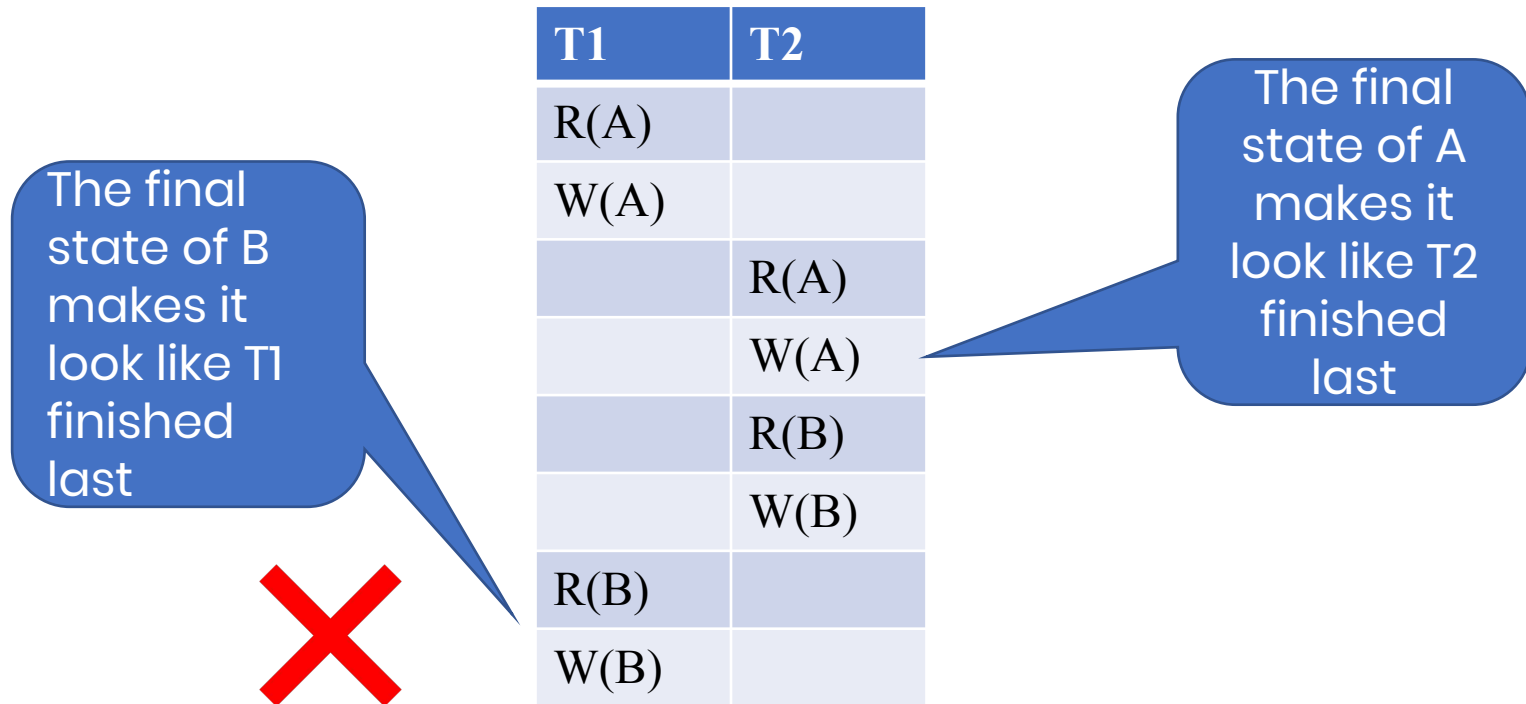


T1	T2
R(A)	
W(A)	
	R(A)
	W(A)
	R(B)
	W(B)
R(B)	
W(B)	

# Serializable Schedules

- Not serializable to either order

$R_1(A), W_1(A), R_2(A), W_2(A), R_2(B), W_2(B), R_1(B), W_1(B)$



# Checking Serializability

- How does the DBMS tell if some schedule is serializable?

# Conflicts

## **Conflict**

A pair of actions in a schedule such that, if their order is interchanged, then the behavior of at least one of the transactions involved can change.

# Conflicts

- In terms of the application concurrency problems we saw:
- Lost Update □ Write-Write (WW) conflict
- Dirty Read □ Write-Read (WR) conflict
- Unrepeatable Read □ Read-Write (RW) conflict
- Phantom Read
  - We'll talk about this later

Individual conflicts aren't "bad"! We expect that different transactions have conflicting actions. The problem comes with interleaving the conflicts.

# Types of Conflicts

- Changing the order of things in conflict will cause program behavior to behave badly
- **Intra-transaction conflicts**
  - Operations within a transaction cannot be swapped (you would be literally changing the program)
- **Inter-transaction conflicts**
  - WW conflicts □  $W1(X), W2(X)$
  - WR conflicts □  $W1(X), R2(X)$
  - RW conflicts □  $R1(X), W2(X)$

# Types of Conflicts

- Changing the order of things in conflict will cause program behavior to behave badly

- **Intra-transaction conflicts**

- Operations within a transaction cannot be swapped (you would be literally changing the program)

- **Inter-transaction conflicts**

- WW conflicts  $\square W1(X), W2(X)$
  - WR conflicts  $\square W1(X), R2(X)$
  - RW conflicts  $\square R1(X), W2(X)$

Note what's missing:  
actions on different elements or reading the same element.  
are NOT conflicts.



# Serial Schedules

- A **serial schedule** is a schedule where each transaction would be executed in some order
- A **serializable schedule** is a schedule that is equivalent to a serial schedule
  - If the schedule were executed and you were given a before and after, you would not be able to tell if the transactions were interleaved
- A **conflict serializable schedule** is a schedule that can be transformed into a serial schedule by performing a series of swaps of adjacent non-conflicting actions

# Conflict Serializable Schedule Example

T1	T2
R(A)	
W(A)	
	R(A)
	W(A)
R(B)	
W(B)	
	R(B)
	W(B)

# Conflict Serializable Schedule Example

T1	T2
R(A)	
W(A)	
	R(A)
	W(A)
R(B)	
W(B)	
	R(B)
	W(B)

# Conflict Serializable Schedule Example

T1	T2
R(A)	
W(A)	
	R(A)
R(B)	
	W(A)
W(B)	
	R(B)
	W(B)

# Conflict Serializable Schedule Example

T1	T2
R(A)	
W(A)	
	R(A)
R(B)	
	W(A)
W(B)	
	R(B)
	W(B)

# Conflict Serializable Schedule Example

T1	T2
R(A)	
W(A)	
R(B)	
	R(A)
	W(A)
W(B)	
	R(B)
	W(B)

# Conflict Serializable Schedule Example

T1	T2
R(A)	
W(A)	
R(B)	
	R(A)
	W(A)
W(B)	
	R(B)
	W(B)

# Conflict Serializable Schedule Example

T1	T2
R(A)	
W(A)	
R(B)	
	R(A)
W(B)	
	W(A)
	R(B)
	W(B)



# Conflict Serializable Schedule Example

T1	T2
R(A)	
W(A)	
R(B)	
	R(A)
W(B)	
	W(A)
	R(B)
	W(B)

# Conflict Serializable Schedule Example

T1	T2
R(A)	
W(A)	
R(B)	
W(B)	
	R(A)
	W(A)
	R(B)
	W(B)

# Non Conflict Serializable Schedule Example

T1	T2
R(A)	
W(A)	
	R(A)
	W(A)
	R(B)
	W(B)
R(B)	
W(B)	

# Non Conflict Serializable Schedule Example

T1	T2
R(A)	
W(A)	
	R(A)
	W(A)
	R(B)
R(B)	
	W(B)
W(B)	



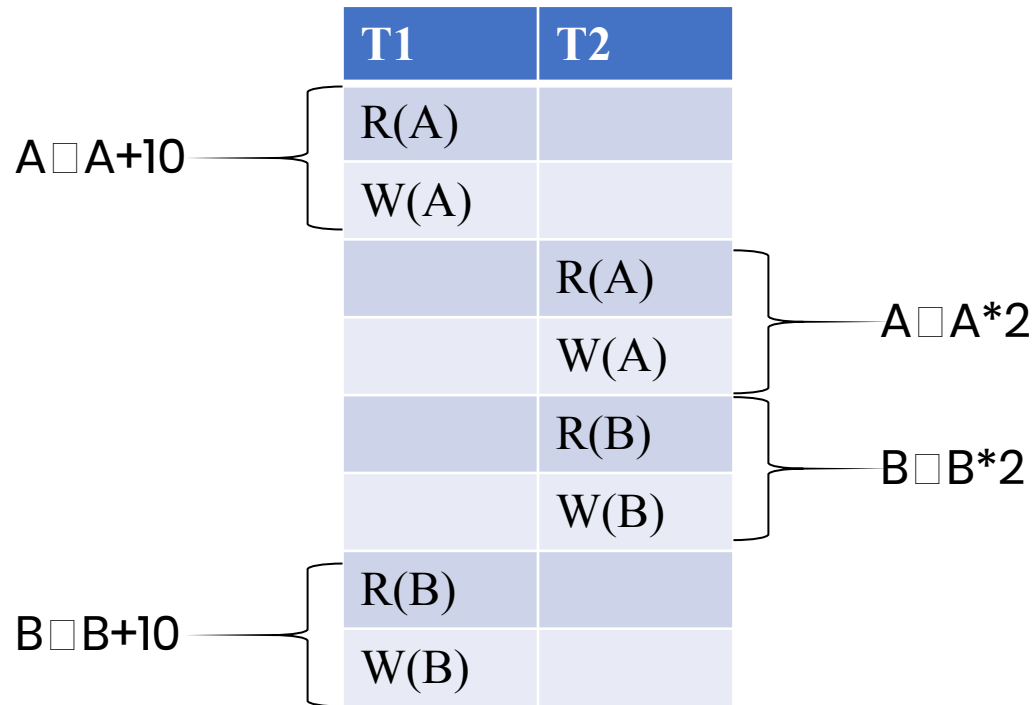
Conflict rule broken!

# Conflict Serializability

- Showing program serializability is hard
  - Needs lots of extra information besides R, W, I, D
- Observation: Enforce something something simpler but stronger than serializability
- **Conflict serializability implies serializability**
- Serializability does not imply conflict serializability

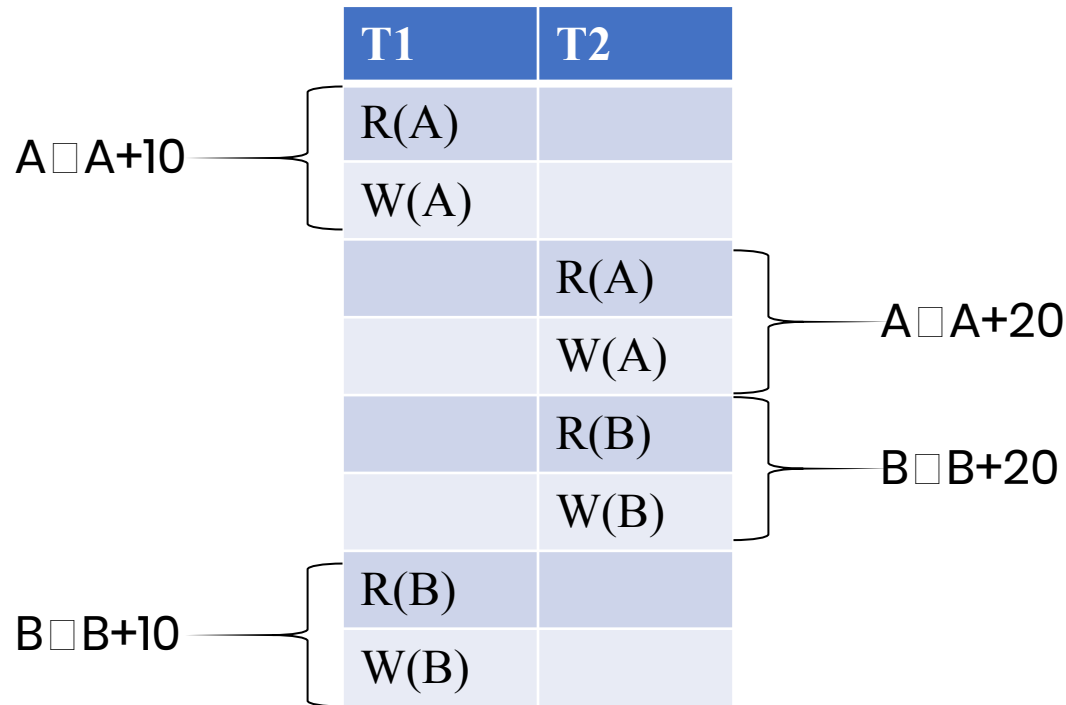
# Serializable vs Conflict Serializable

Not serializable nor conflict serializable



# Serializable vs Conflict Serializable

Serializable but not conflict serializable






# Enforcing Conflict Serializability

- We only care if some conflict rule would be broken
- But we need an efficient algorithm
- Method:
  - Model each transaction as a node
  - Model a inter-transaction conflict as a directed edge
  - If the resulting graph is a DAG then there is a serial order
  - Conflict serializability enforcement turns into the graph cycle detection problem




# Example of Checking Conflict Serializability

$R_1(A), W_2(A), W_2(C), R_3(A), W_3(A), R_3(B), W_1(B)$




# Example of Checking Conflict Serializability

$R_1(A), W_2(A), W_2(C), R_3(A), W_3(A), R_3(B), W_1(B)$




# Example of Checking Conflict Serializability

$R_1(A), W_2(A), W_2(C), R_3(A), W_3(A), R_3(B), W_1(B)$




# Example of Checking Conflict Serializability

$R_1(A), W_2(A), W_2(C), R_3(A), W_3(A), R_3(B), W_1(B)$



# Example of Checking Conflict Serializability




$R_1(A), W_2(A), W_2(C), R_3(A), W_3(A), R_3(B), W_1(B)$



Cycle ☐ Not conflict serializable


# Example of Checking Conflict Serializability

$R_1(X), W_2(X), W_1(Y), W_2(Y), W_3(Y), R_3(X), W_3(X)$




# Example of Checking Conflict Serializability

$R_1(X), W_2(X), W_1(Y), W_2(Y), W_3(Y), R_3(X), W_3(X)$



# Example of Checking Conflict Serializability


$R_1(X)$ ,  $W_2(X)$ ,  $W_1(Y)$ ,  $W_2(Y)$ ,  $W_3(Y)$ ,  $R_3(X)$ ,  $W_3(X)$






# Example of Checking Conflict Serializability

$R_1(X)$ ,  $W_2(X)$ ,  $W_1(Y)$ ,  $W_2(Y)$ ,  $W_3(Y)$ ,  $R_3(X)$ ,  $W_3(X)$



# Example of Checking Conflict Serializability

$R_1(X), W_2(X), W_1(Y), W_2(Y), W_3(Y), R_3(X), W_3(X)$



DAG ☐ Conflict serializable ☐ Serializable

# Takeaways

- When a database has multiple concurrent users, a variety of conflicts can happen.
- DBMSs implement transactions with ACID properties to avoid anomalies while allowing concurrent users.