

Introduction to Data Management

NoSQL

Alyssa Pittman

Based on slides by Jonathan Leang, Dan Suciu, et al

Paul G. Allen School of Computer Science and Engineering
University of Washington, Seattle

Classical Database Application Problems

OLTP
(Online Transaction Processing)

OLAP
(Online Analytical Processing)

Classical Database Application Problems

OLTP (Online Transaction Processing)	OLAP (Online Analytical Processing)
Transaction-heavy workloads	Complex query workloads

Classical Database Application Problems

OLTP (Online Transaction Processing)	OLAP (Online Analytical Processing)
Transaction-heavy workloads	Complex query workloads
Many simple lookup or single-join queries	Many joins, aggregations, etc.

Classical Database Application Problems

OLTP (Online Transaction Processing)	OLAP (Online Analytical Processing)
Transaction-heavy workloads	Complex query workloads
Many simple lookup or single-join queries	Many joins, aggregations, etc.
Many small updates and inserts	Little to no updates

Classical Database Application Problems

OLTP (Online Transaction Processing)	OLAP (Online Analytical Processing)
Transaction-heavy workloads	Complex query workloads
Many simple lookup or single-join queries	Many joins, aggregations, etc.
Many small updates and inserts	Little to no updates
Managing consistency is critical	Query optimization and processing is critical

Classical Database Application Problems

OLTP (Online Transaction Processing)	OLAP (Online Analytical Processing)
Transaction-heavy workloads	Complex query workloads
Many simple lookup or single-join queries	Many joins, aggregations, etc.
Many small updates and inserts	Little to no updates
Managing consistency is critical	Query optimization and processing is critical
Flights, banking, etc. (many users)	Business intelligence (few users)

Client-Server Applications



Client-Server Applications

Single server runs
the entire
database



Client-Server Applications

Single server runs
the entire
database



Could be:

- Your own computer
- Cloud-hosted DB



Client-Server Applications

Single server runs the entire database



Could be:

- Your own computer
- Cloud-hosted DB

Multiple client applications connect to DB server

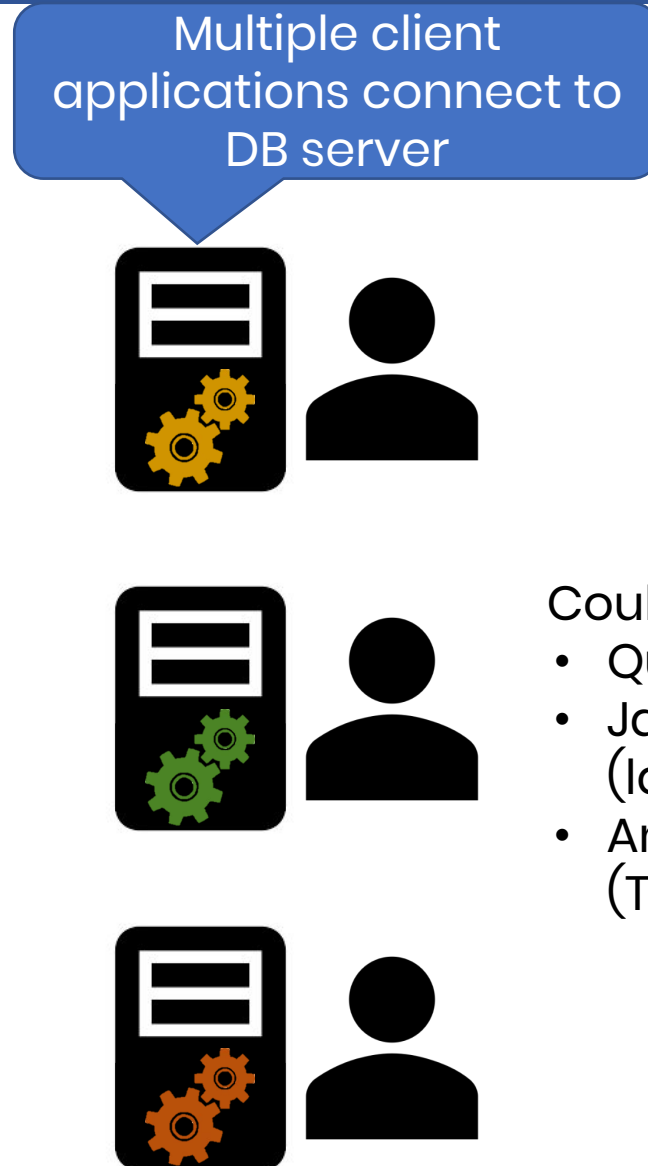


Client-Server Applications



Could be:

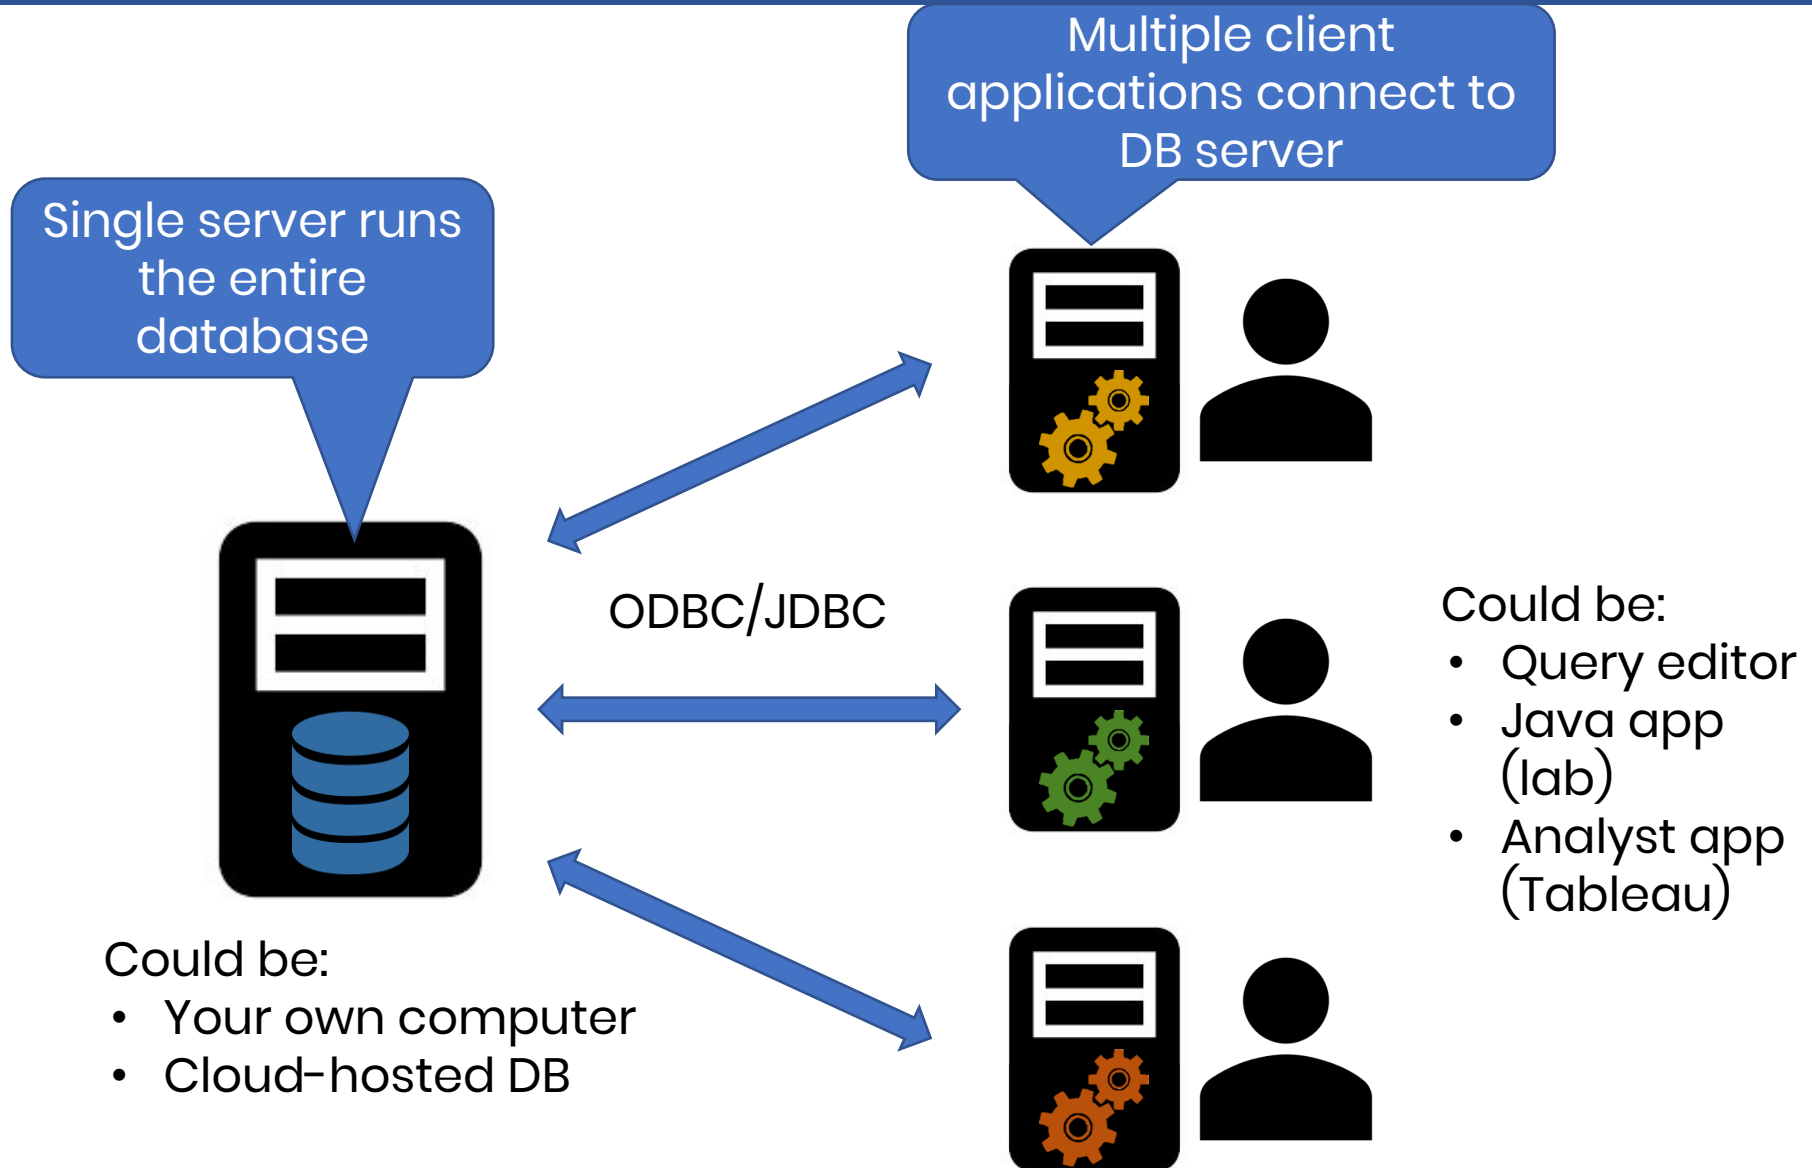
- Your own computer
- Cloud-hosted DB



Could be:

- Query editor
- Java app (Ib)
- Analyst app (Tableau)

Client-Server Applications



Client-Server Applications

Sufficient for OLAP (simple)

Can't scale connections for OLTP

Single server runs the entire database



ODBC/JDBC

Multiple client applications connect to DB server



Could be:

- Your own computer
- Cloud-hosted DB

Could be:

- Query editor
- Java app (Ib)
- Analyst app (Tableau)

The World Wide Web – Web 2.0

- A new class of problem emerges in the late 90s and early 2000s (and is still a problem today)
- What is Web 2.0?
 - Social web (Facebook, Amazon, Instagram, ...)
 - Startup services need to **scale quickly by orders of magnitude** (shared-nothing architecture!)
 - **Exclusively OLTP workloads**

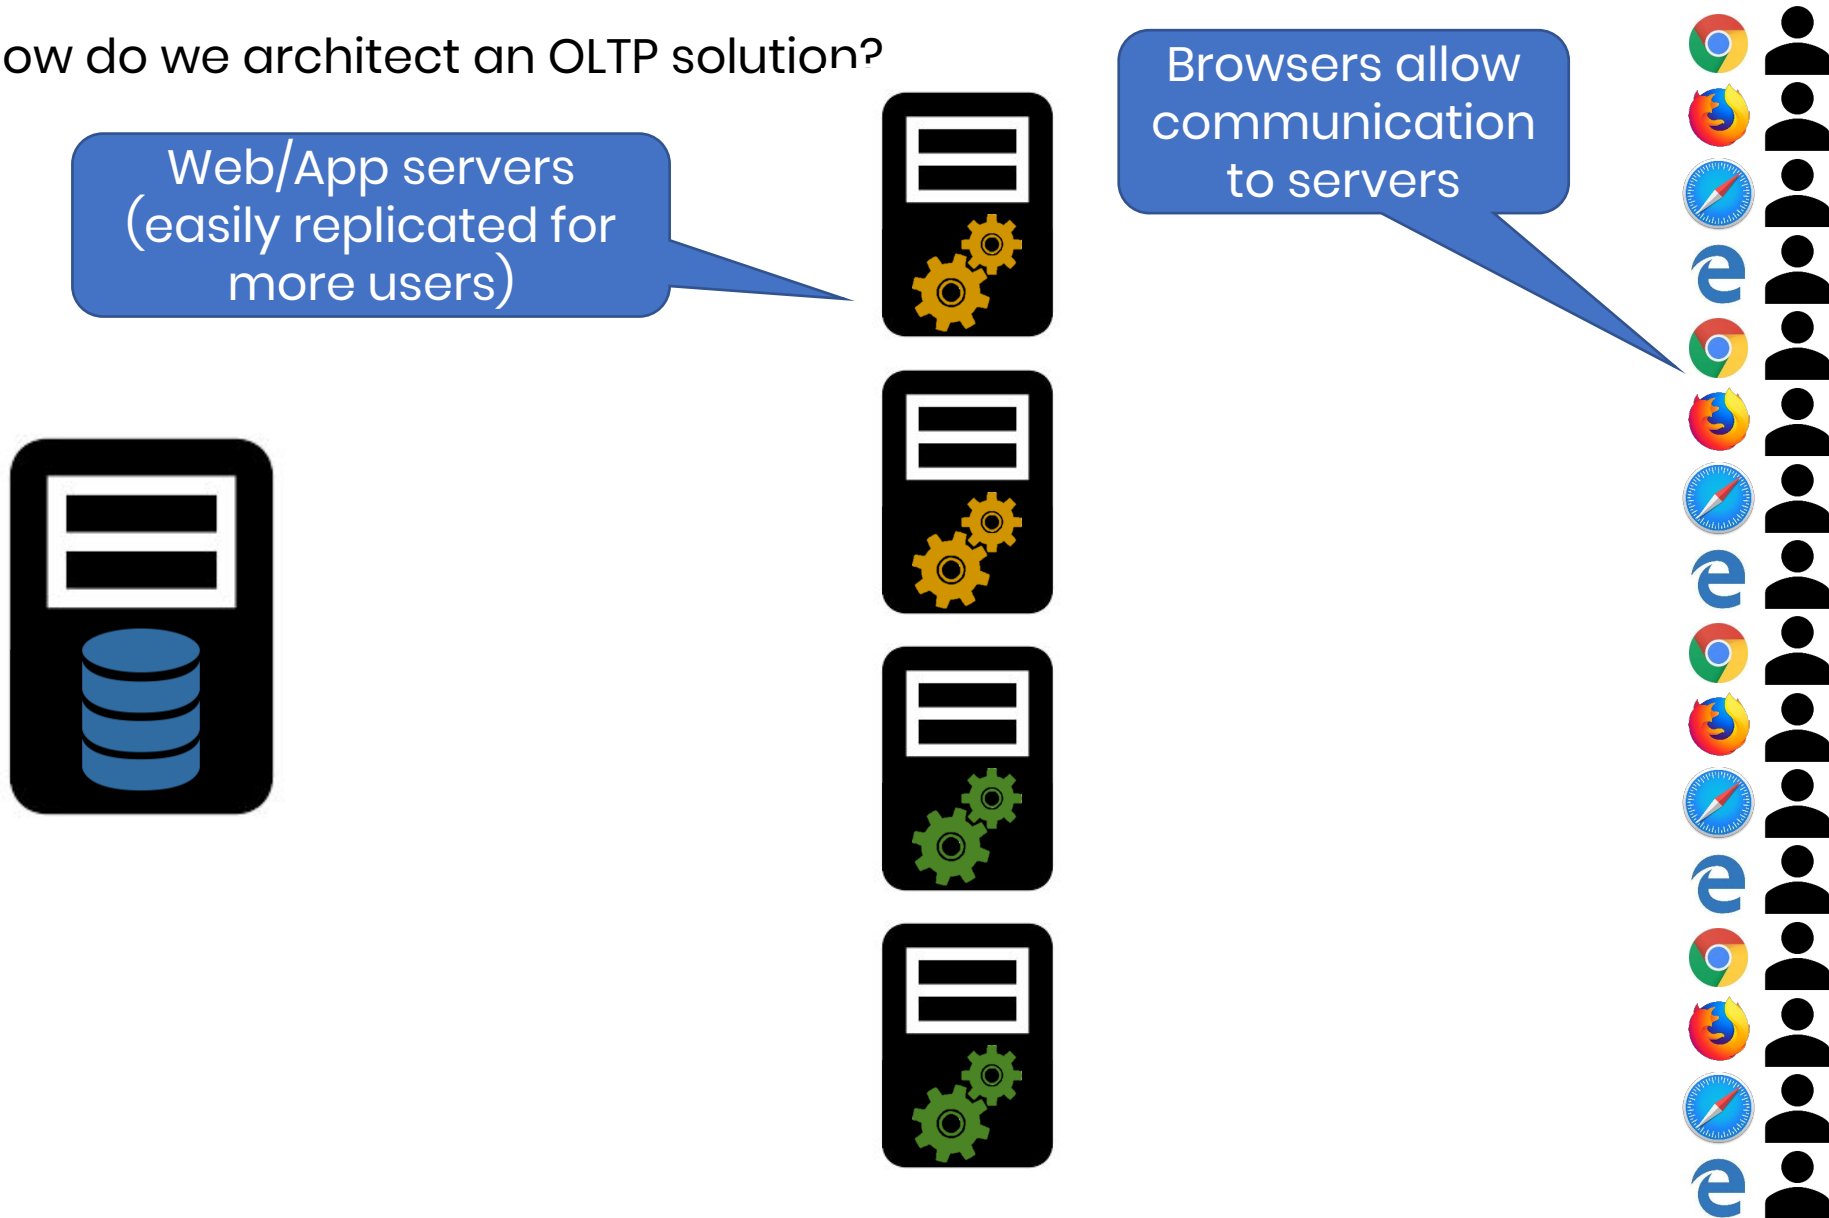
3-Tiered Web Architecture

How do we architect an OLTP solution?



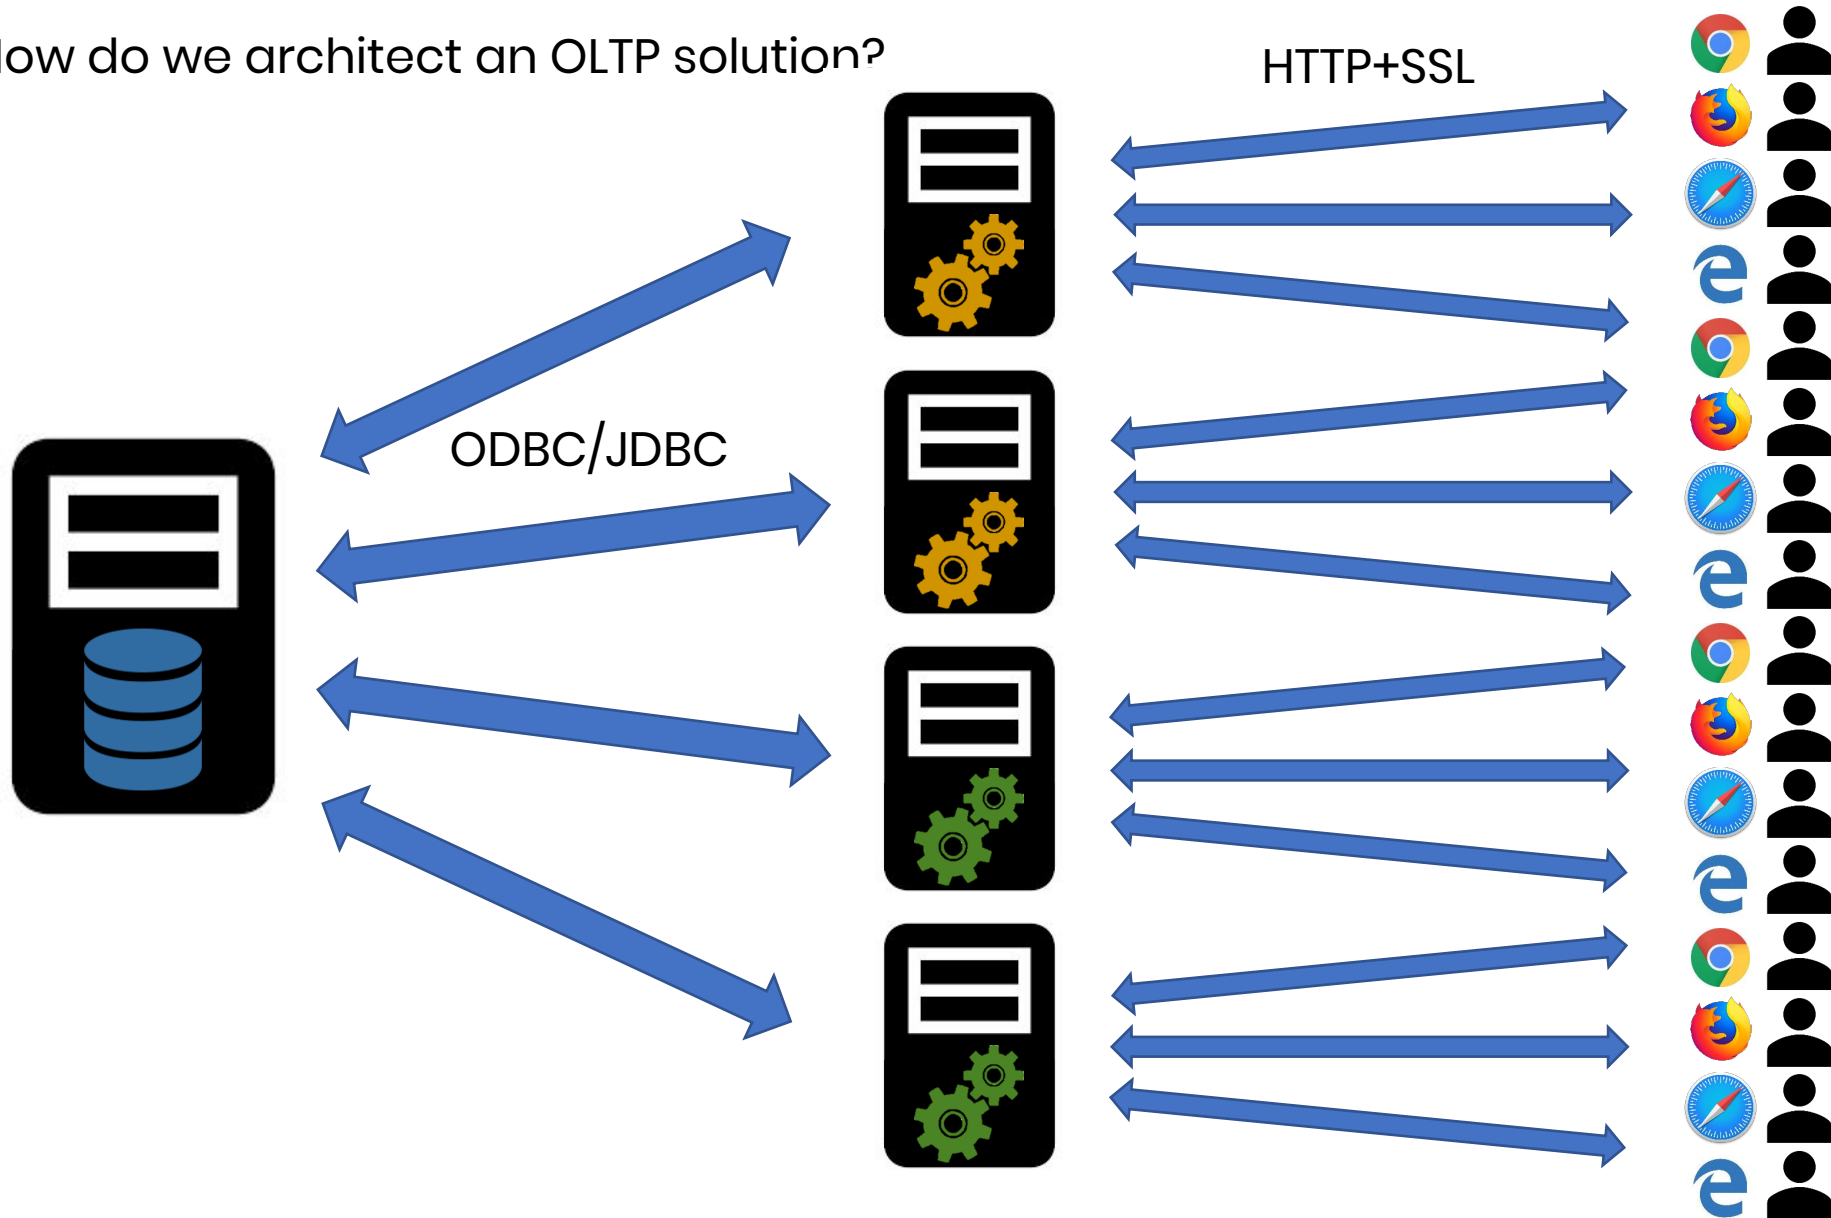
3-Tiered Web Architecture

How do we architect an OLTP solution?



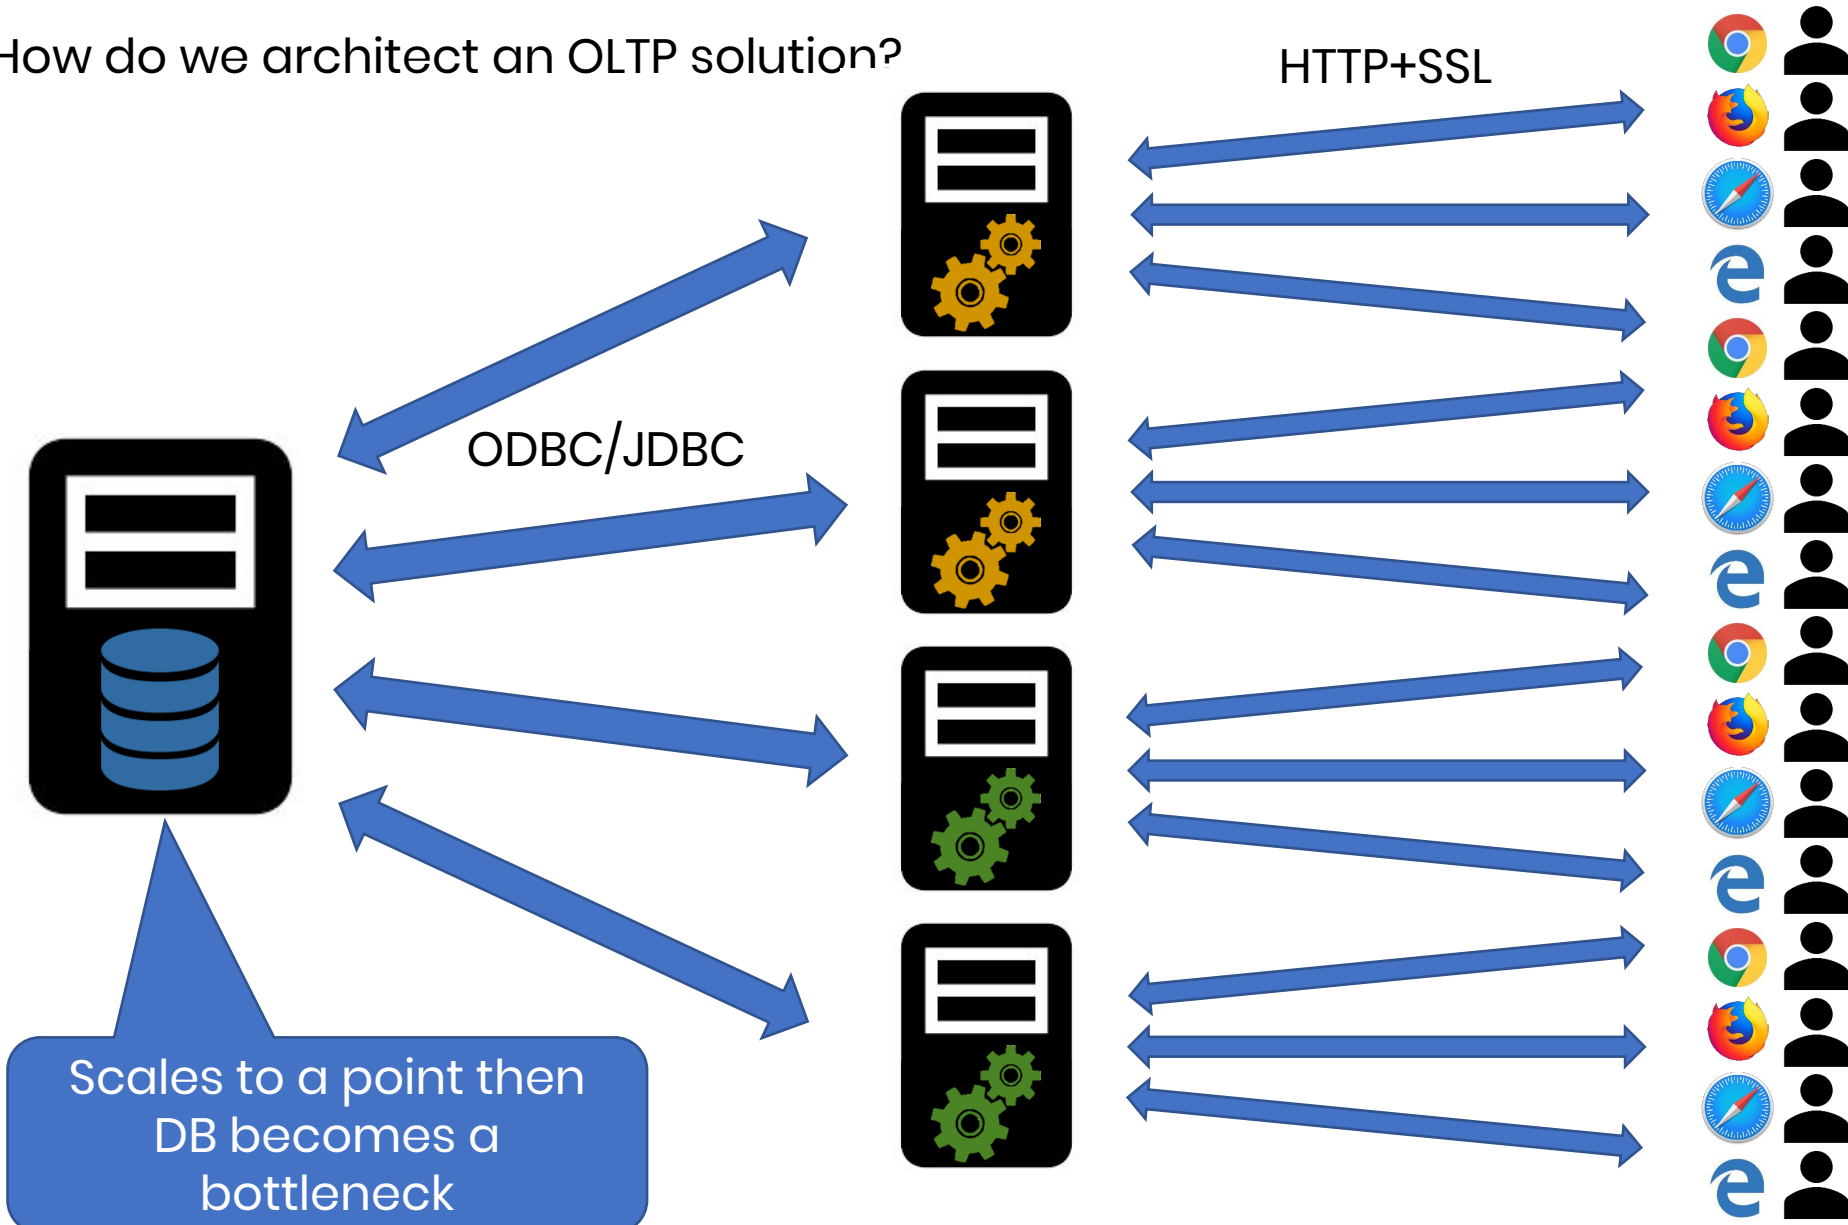
3-Tiered Web Architecture

How do we architect an OLTP solution?



3-Tiered Web Architecture

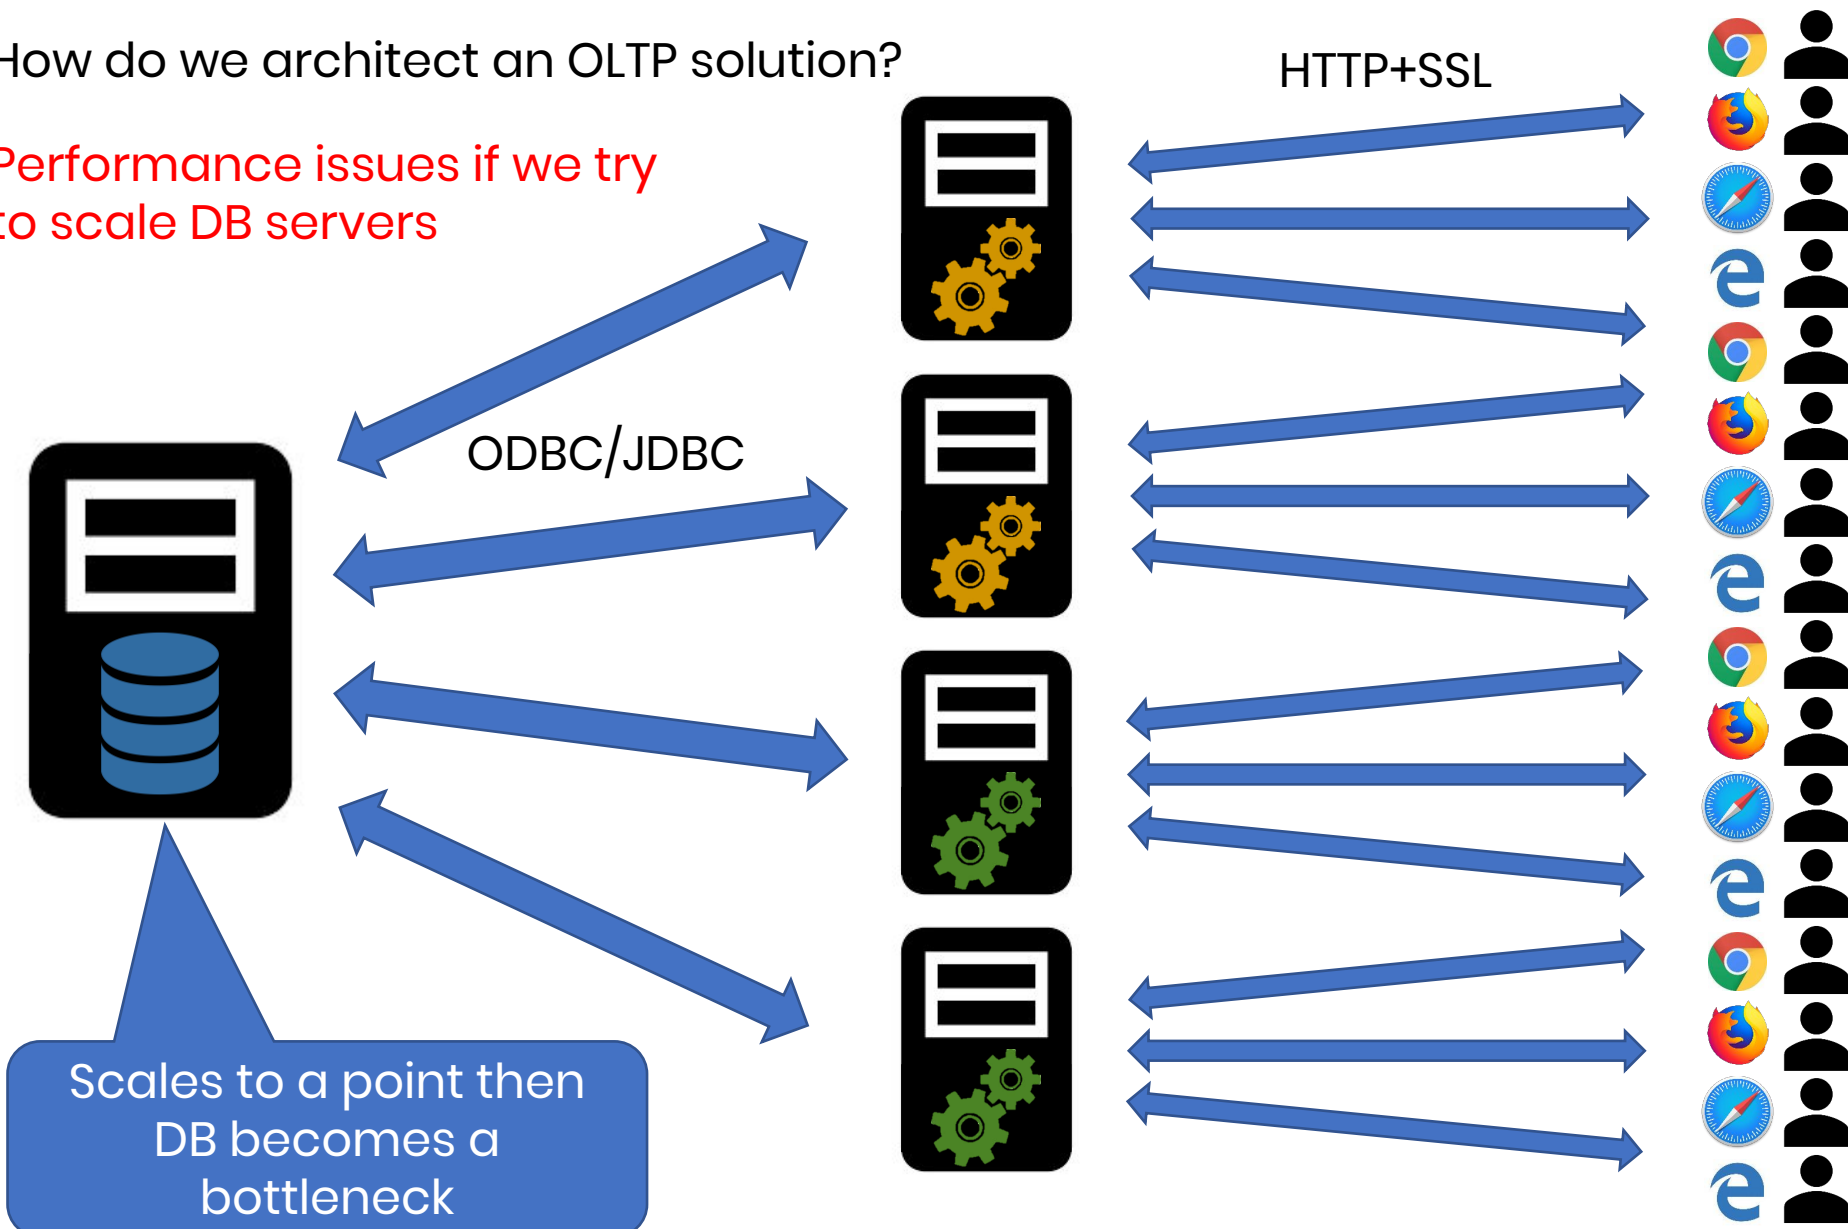
How do we architect an OLTP solution?



3-Tiered Web Architecture

How do we architect an OLTP solution?

Performance issues if we try
to scale DB servers



Database Scaling Techniques

- Scale up via:
 - **Partitioning** (sharding)
 - **Replication**

RDBMS Partitioning

- Use multiple machines to distribute data
- Write performance ok
- Read performance suffers!
 - Join across servers may have huge network IO cost

RDBMS Replication

- Create multiple copies of each database partition
- Improves fault tolerance
- Read performance ok
- Write performance suffers!
 - Need to write same value to multiple servers

Distributed RDBMS Consistency Bottleneck

- RDBMS scaling makes consistency hard
 - Partitioning: Need to coordinate server actions
 - Replication: Need to prevent inconsistent versions
 - ACID is hard to maintain

A hashtag on Twitter for a [meetup](#) in San Francisco to discuss systems like Google BigTable, Amazon Dynamo, CouchDB, etc.

Event Details

Introduction

This meetup is about "open source, distributed, non relational databases".

Have you run into limitations with traditional relational databases? Don't mind trading a query language for scalability? Or perhaps you just like shiny new things to try out? Either way this meetup is for you.

Join us in figuring out why these newfangled Dynamo clones and BigTables have become so popular lately. We have gathered presenters from the most interesting projects around to give us all an introduction to the field.

Preliminary schedule

09.45: Doors open

10.00: **Intro session** (Todd Lipcon, Cloudera)

10.40: **Voldemort** (Jay Kreps, LinkedIn)

11.20: Short break

11.30: **Cassandra** (Avinash Lakshman, Facebook)

12.10: Free lunch (sponsored by Last.fm)

13.10: **Dynomite** (Cliff Moon, Powerset)

13.50: **HBase** (Ryan Rawson, Stumbleupon)

14.30: Short break

14.40: **Hypertable** (Doug Judd, Zvents)

15.20: **CouchDB** (Chris Anderson, couch.io)

16.00: Short break

16.10: Lightning talks

16.40: Panel discussion

17.00: Relocate to Kate O'Brien's, 579 Howard St. @ 2nd. First round sponsored by Digg

Registration

The event is free but space is limited, please register if you wish to attend.

Location

Magma room, CBS interactive

235 Second Street

San Francisco, CA 94105



thruddb @thruddb · 23 May 2009

sucks i'm not on the west coast and will not be able to attend [#nosql](#)



Todd Lipcon @tliipcon · 23 May 2009

working on slides for the [#nosql](#) meetup in June. trying to cover all of dist systems in 40 minutes is not as easy as it sounds.



Chris Anderson @jchris · 13 May 2009

It's official - I'll be relaxifying everyone with CouchDB at [#NoSQL](#). Thanks @skr!



A hashtag on Twitter for a [meetup](#) in San Francisco to discuss systems like Google BigTable, Amazon Dynamo, CouchDB, etc.

Because #NoRDBMS doesn't have quite the same ring to it

Event Details

Introduction

This meetup is about "open source, distributed, non relational databases".

Have you run into limitations with traditional relational databases? Don't mind trading a query language for scalability? Or perhaps you just like shiny new things to try out? Either way this meetup is for you.

Join us in figuring out why these newfangled Dynamo clones and BigTables have become so popular lately. We have gathered presenters from the most interesting projects around to give us all an introduction to the field.

Preliminary schedule

09.45: Doors open

10.00: **Intro session** (Todd Lipcon, Cloudera)

10.40: **Voldemort** (Jay Kreps, LinkedIn)

11.20: Short break

11.30: **Cassandra** (Avinash Lakshman, Facebook)

12.10: Free lunch (sponsored by Last.fm)

13.10: **Dynomite** (Cliff Moon, Powerset)

13.50: **HBase** (Ryan Rawson, Stumbleupon)

14.30: Short break

14.40: **Hypertable** (Doug Judd, Zvents)

15.20: **CouchDB** (Chris Anderson, couch.io)

16.00: Short break

16.10: Lightning talks

16.40: Panel discussion

17.00: Relocate to Kate O'Brien's, 579 Howard St. @ 2nd. First round sponsored by Digg

Registration

The event is free but space is limited, please register if you wish to attend.

Location

Magma room, CBS interactive

235 Second Street

San Francisco, CA 94105



thruodb @thruodb · 23 May 2009

sucks i'm not on the west coast and will not be able to attend [#nosql](#)



Todd Lipcon @tlipcon · 23 May 2009

working on slides for the [#nosql](#) meetup in June. trying to cover all of dist systems in 40 minutes is not as easy as it sounds.



Chris Anderson @jchris · 13 May 2009

It's official - I'll be relaxifying everyone with CouchDB at [#NoSQL](#). Thanks @skr!



How NoSQL Solves Web Scaling



How NoSQL Solves Web Scaling

i give up

NoSQL in a Nutshell

- NoSQL □ Looser data model
 - Give up built-in OLAP/analysis functionality
 - Give up built-in ACID consistency

NoSQL in a Nutshell

- NoSQL works for Web 2.0 business models
 - **No OLAP anyway**
 - **Availability is more important than consistency for Web 2.0**
 - Facebook:
 - I don't care if I don't see every like in real time
 - I care if I can't send a like
 - Amazon:
 - I don't care if my cart forgot an item
 - I care if I can't put an item into my cart

Let's Drop ACID

- RDBMSs have the ACID consistency model
- NoSQL sys. have the **BASE** consistency model
- **Basically Available**
 - Most failures do not cause a complete system outage
- **Soft state**
 - System is not always write-consistent
- **Eventually consistent**
 - Data will eventually converge to agreed values

Why the Sacrifice?

Why can't we have both Consistency and Availability?

CAP Theorem

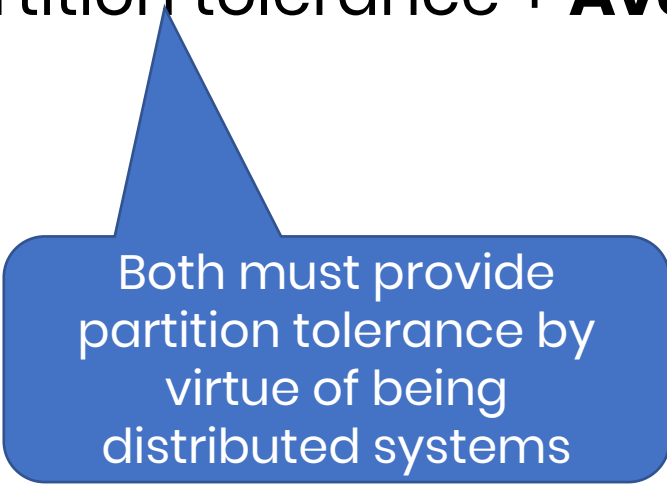
- Old name: Brewer's Conjecture
- In a distributed data store, one can only provide two of the following three guarantees:
 - **Consistency**
 - Every read receives the most recent write or an error
 - **Availability**
 - Every request must respond with a non-error
 - **Partition tolerance**
 - Continued operation in presence of dropped or delayed messages

RDBMS vs NoSQL Systems

- Distributed RDBMS
 - Partition tolerance + **Consistency**
- NoSQL Systems
 - Partition tolerance + **Availability**

RDBMS vs NoSQL Systems

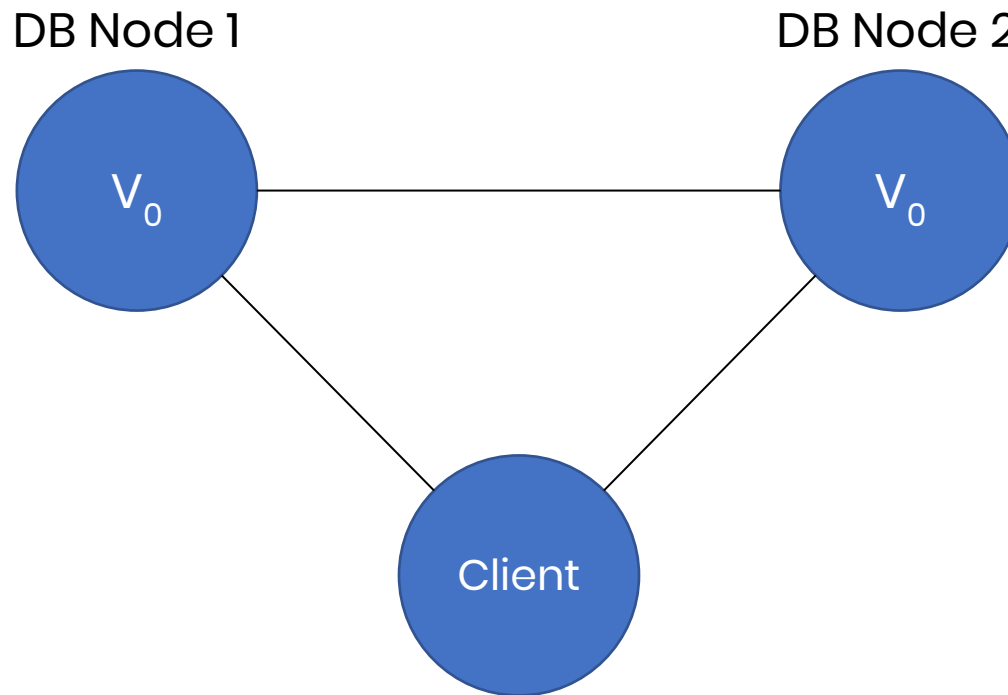
- Distributed RDBMS
 - Partition tolerance + **Consistency**
- NoSQL Systems
 - Partition tolerance + **Availability**



Both must provide
partition tolerance by
virtue of being
distributed systems

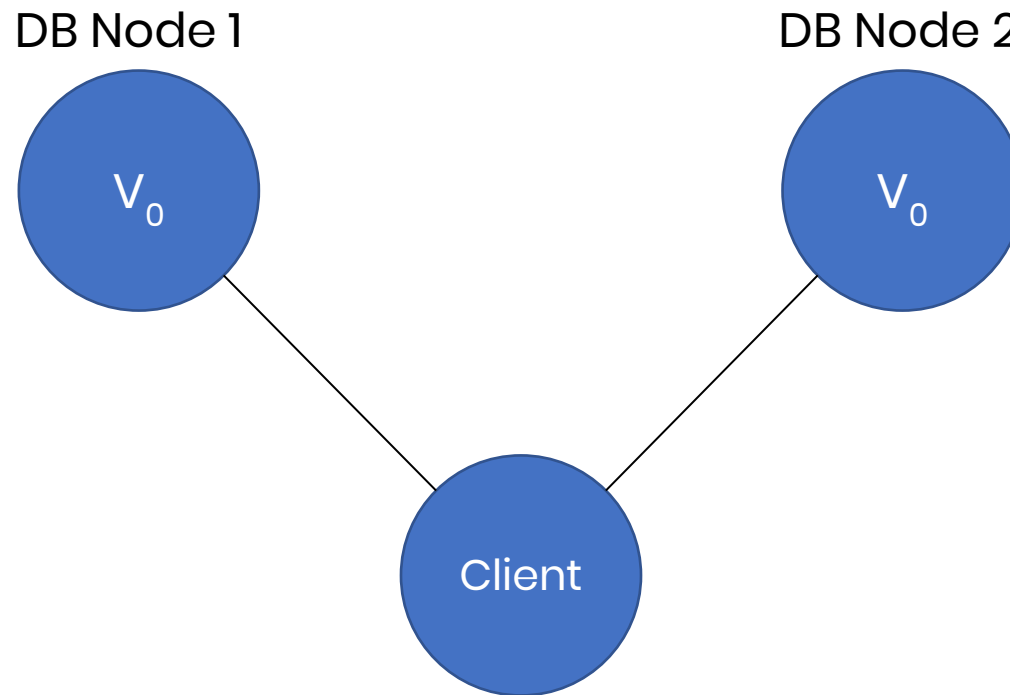
RDBMS vs NoSQL Systems

Partition tolerance + **Consistency**



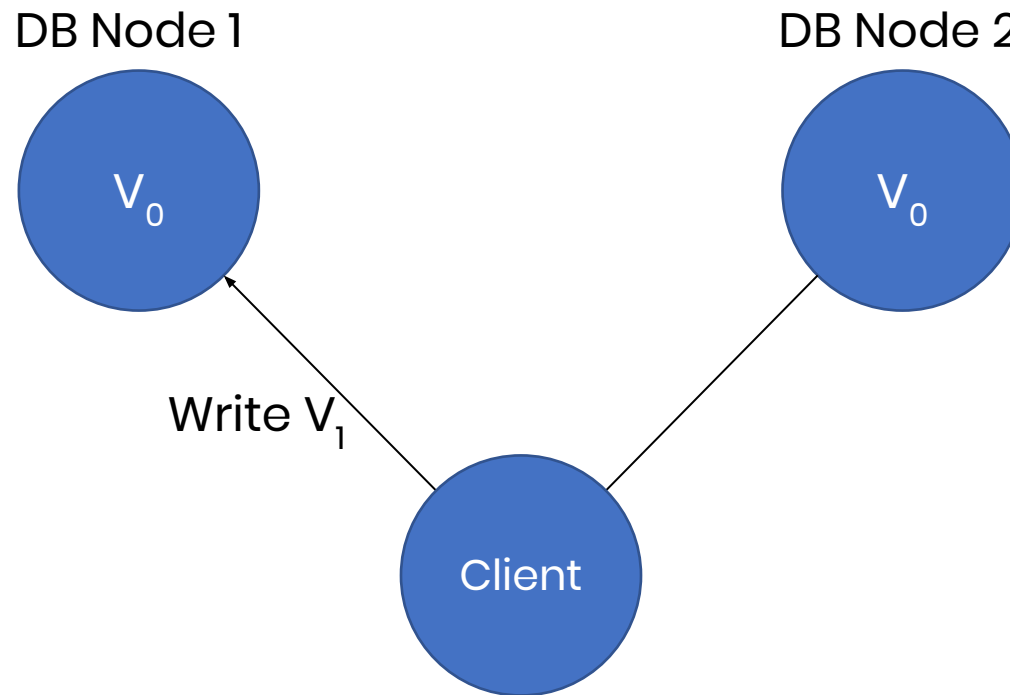
RDBMS vs NoSQL Systems

Partition tolerance + **Consistency**



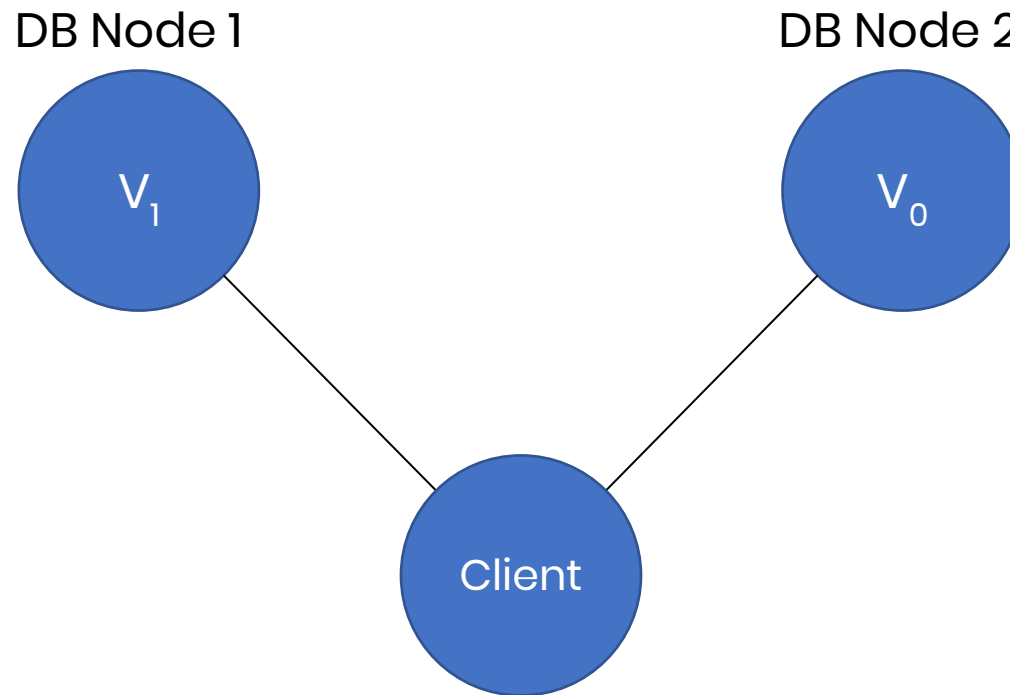
RDBMS vs NoSQL Systems

Partition tolerance + **Consistency**



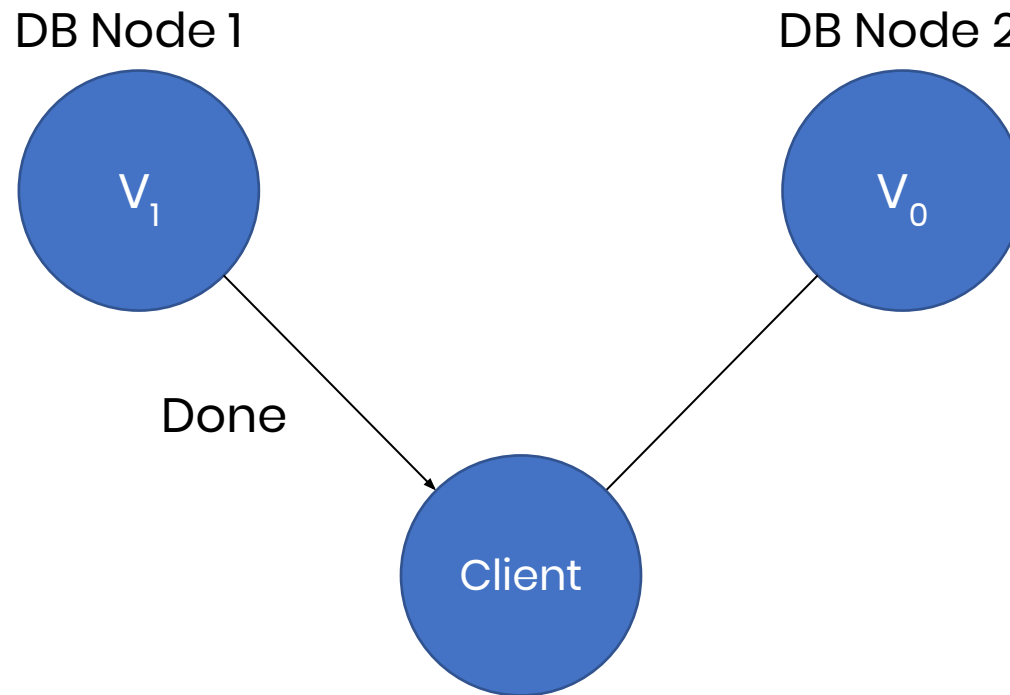
RDBMS vs NoSQL Systems

Partition tolerance + **Consistency**



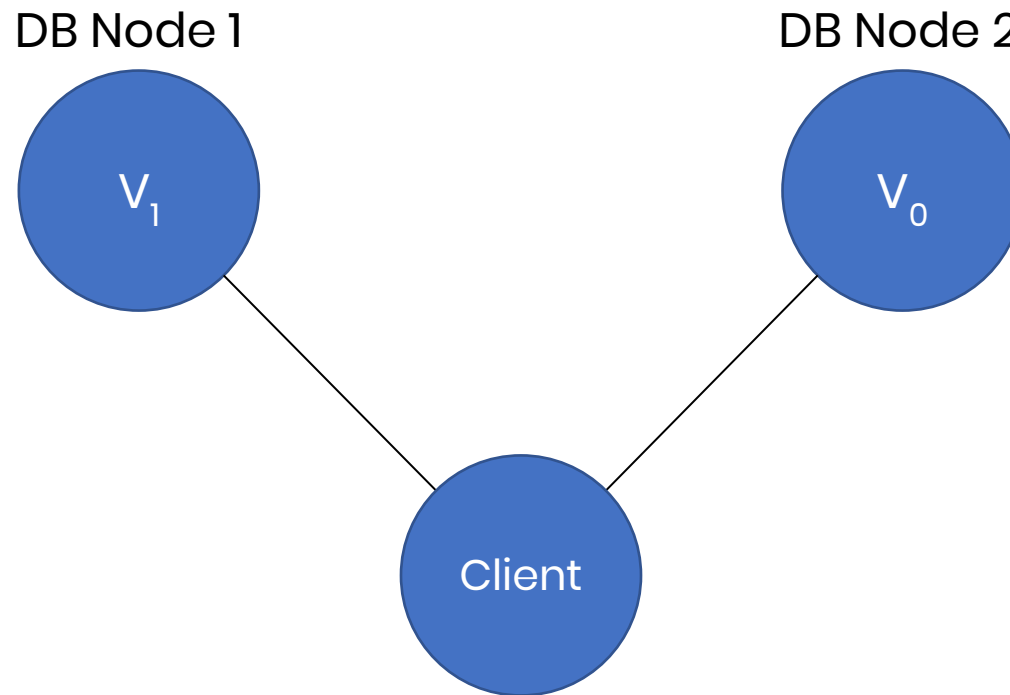
RDBMS vs NoSQL Systems

Partition tolerance + **Consistency**



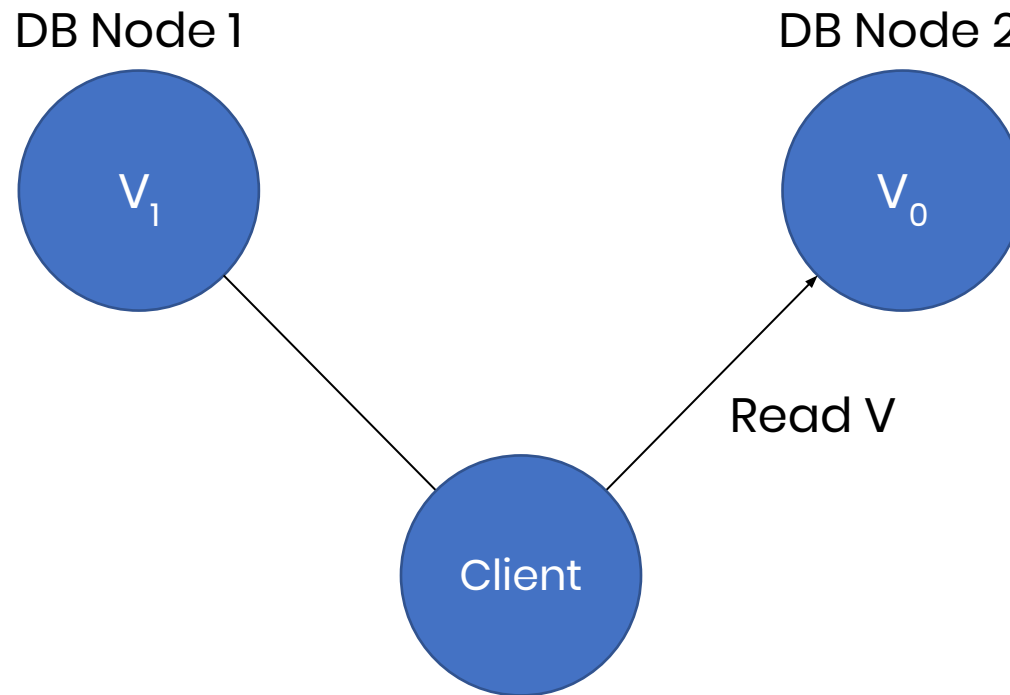
RDBMS vs NoSQL Systems

Partition tolerance + **Consistency**



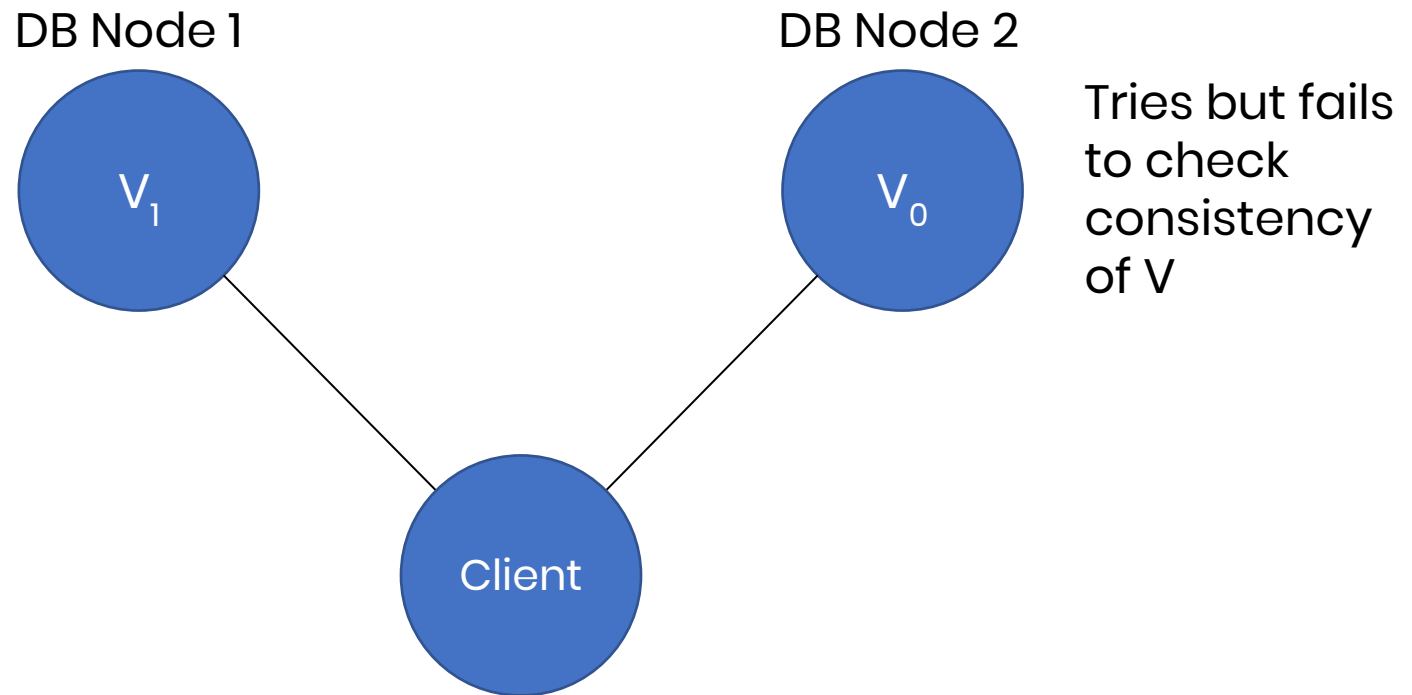
RDBMS vs NoSQL Systems

Partition tolerance + **Consistency**



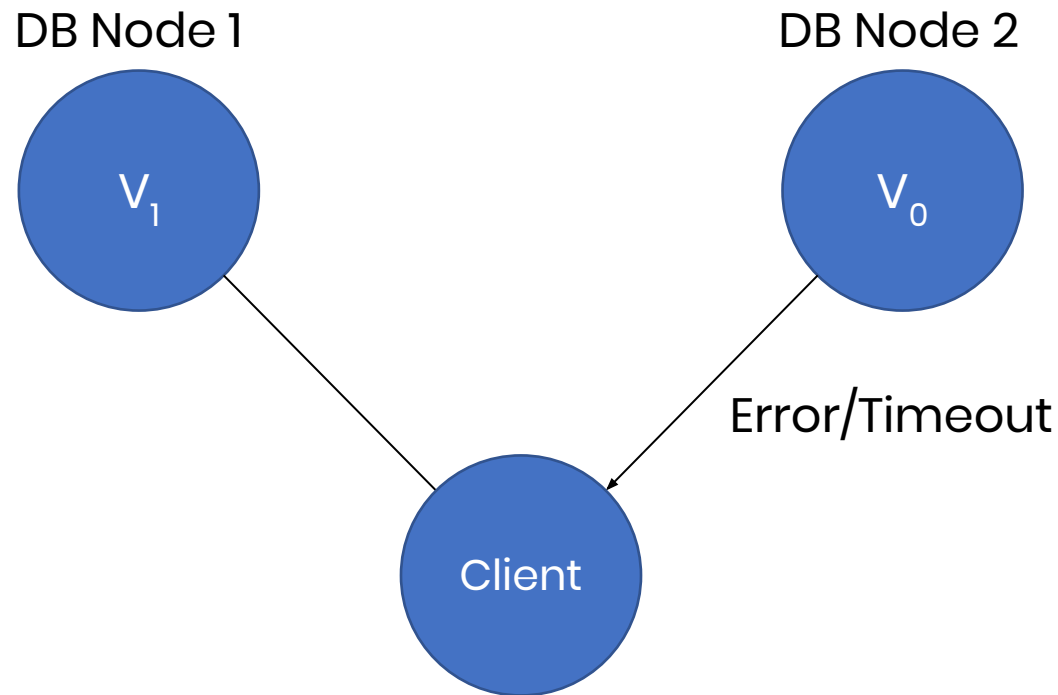
RDBMS vs NoSQL Systems

Partition tolerance + **Consistency**



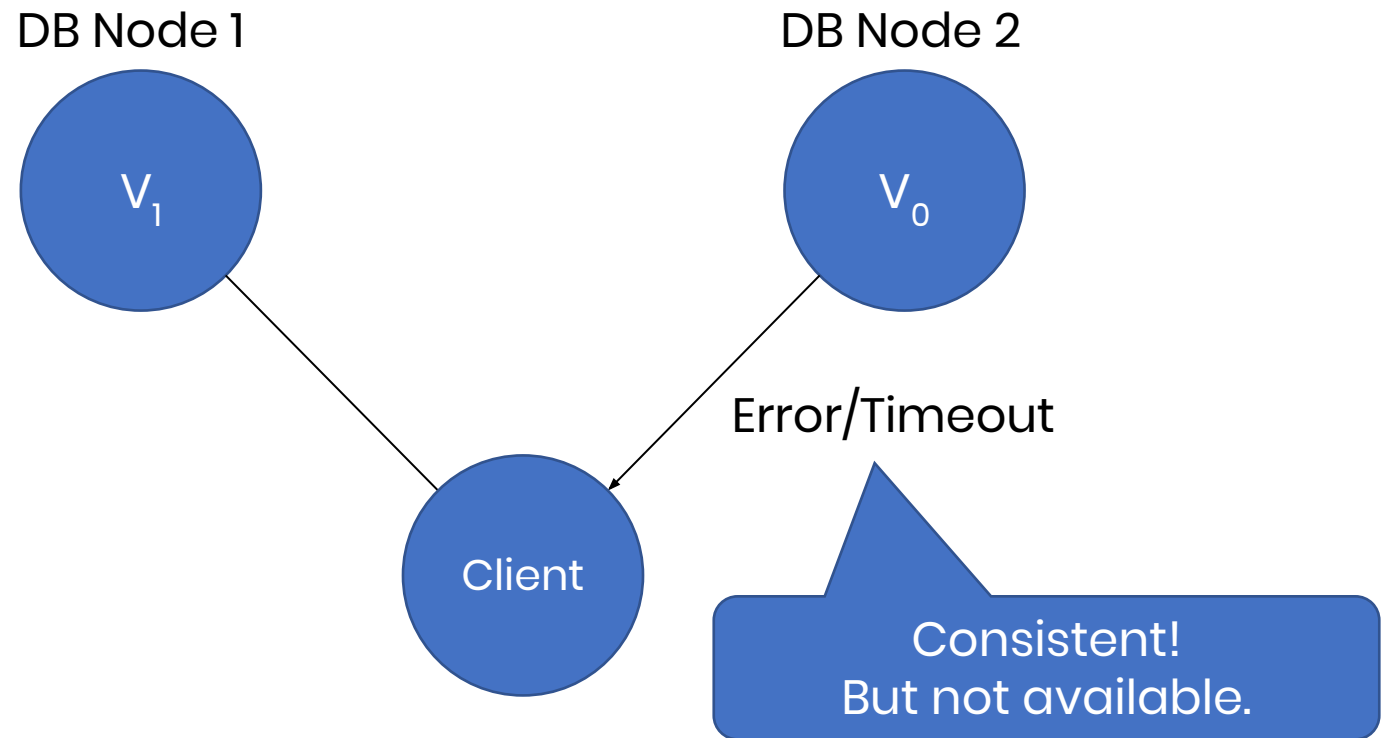
RDBMS vs NoSQL Systems

Partition tolerance + **Consistency**



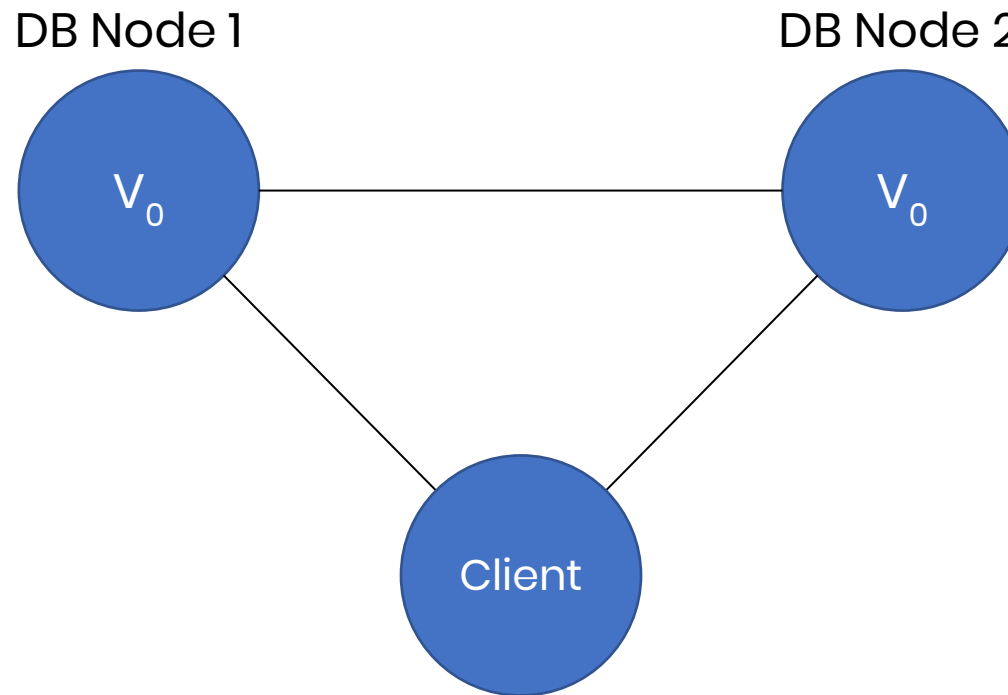
RDBMS vs NoSQL Systems

Partition tolerance + **Consistency**



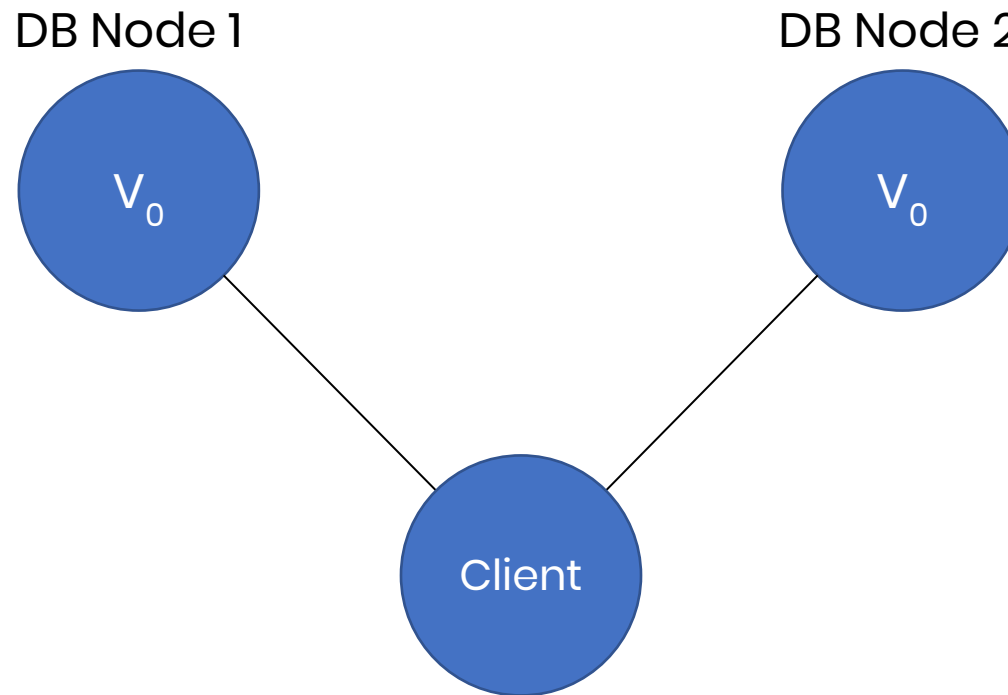
RDBMS vs NoSQL Systems

Partition tolerance + **Availability**



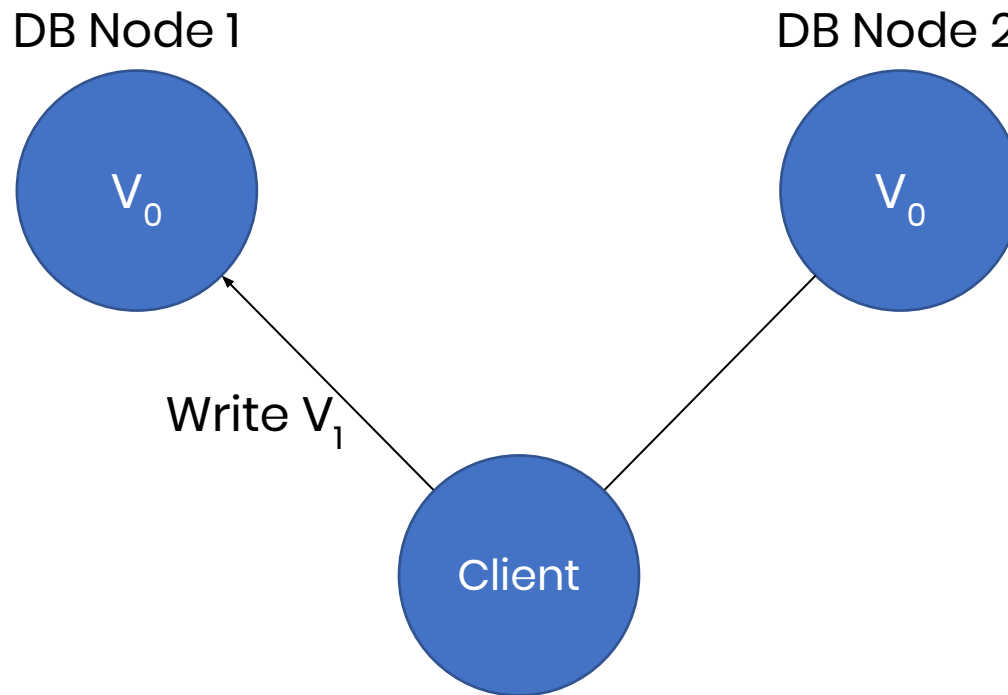
RDBMS vs NoSQL Systems

Partition tolerance + **Availability**



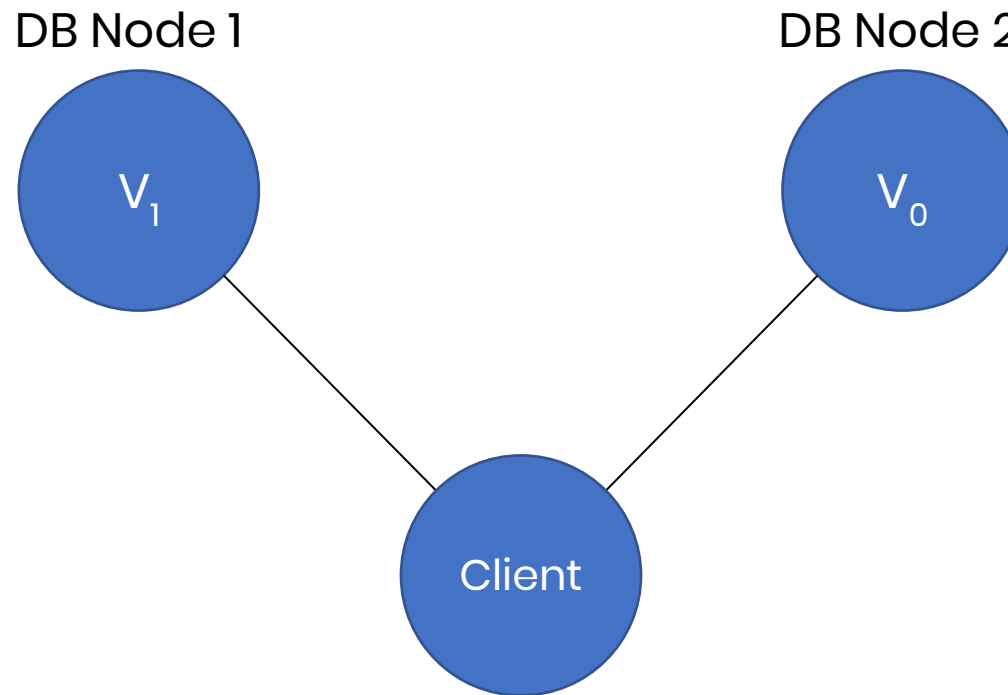
RDBMS vs NoSQL Systems

Partition tolerance + **Availability**



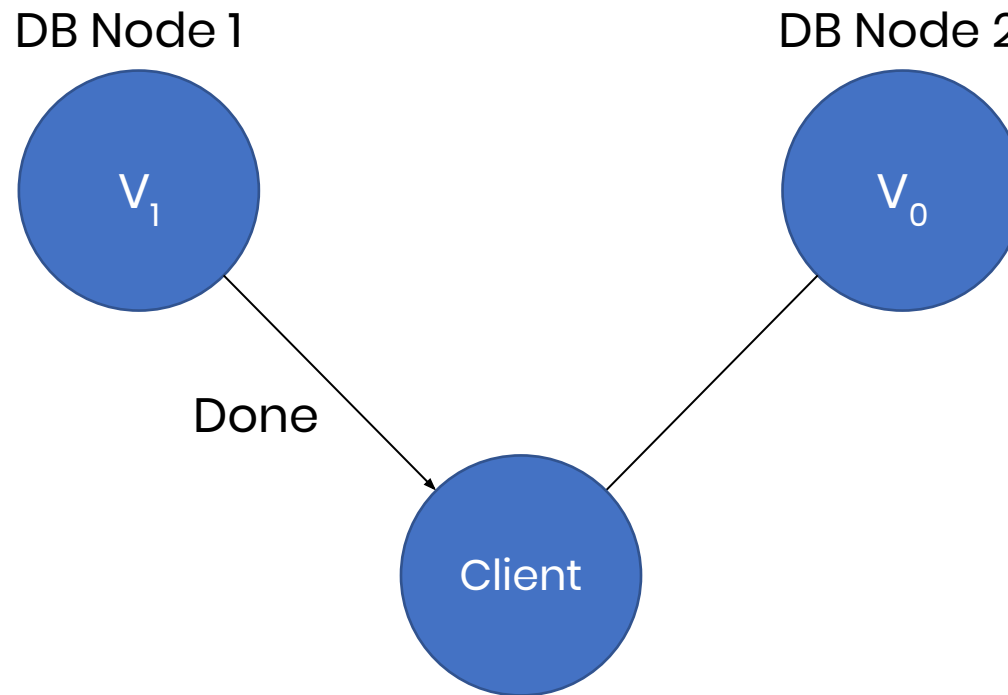
RDBMS vs NoSQL Systems

Partition tolerance + **Availability**



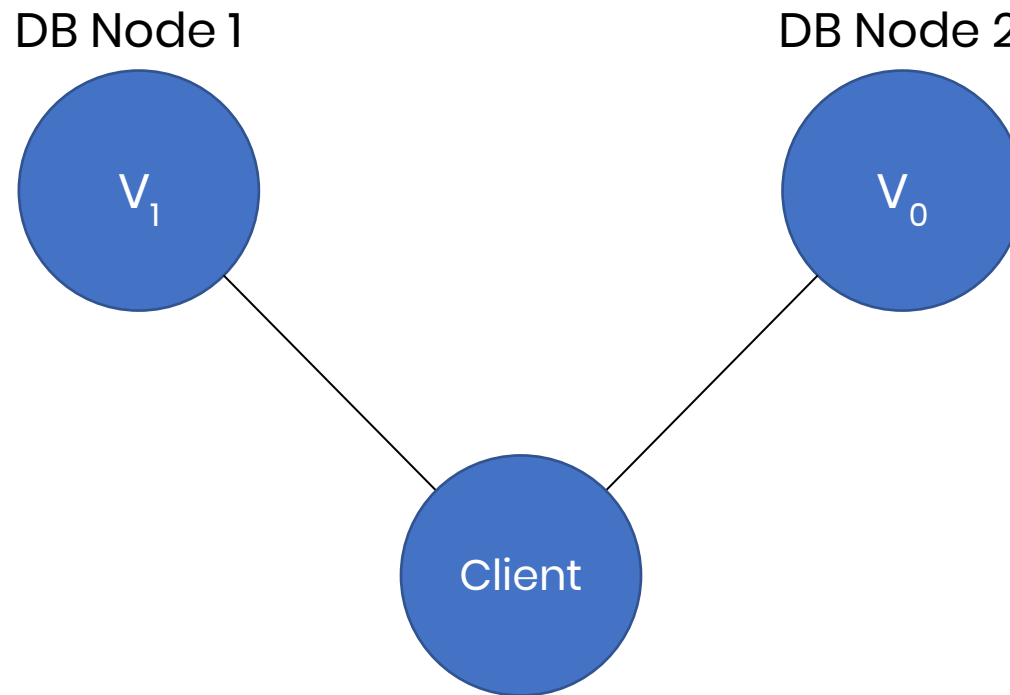
RDBMS vs NoSQL Systems

Partition tolerance + **Availability**



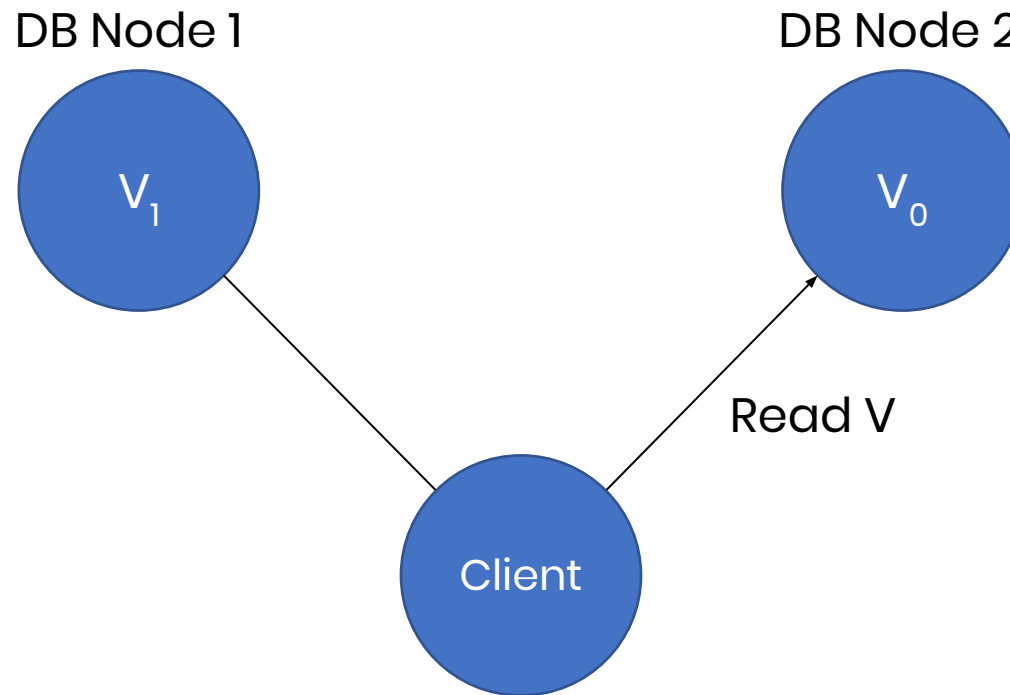
RDBMS vs NoSQL Systems

Partition tolerance + **Availability**



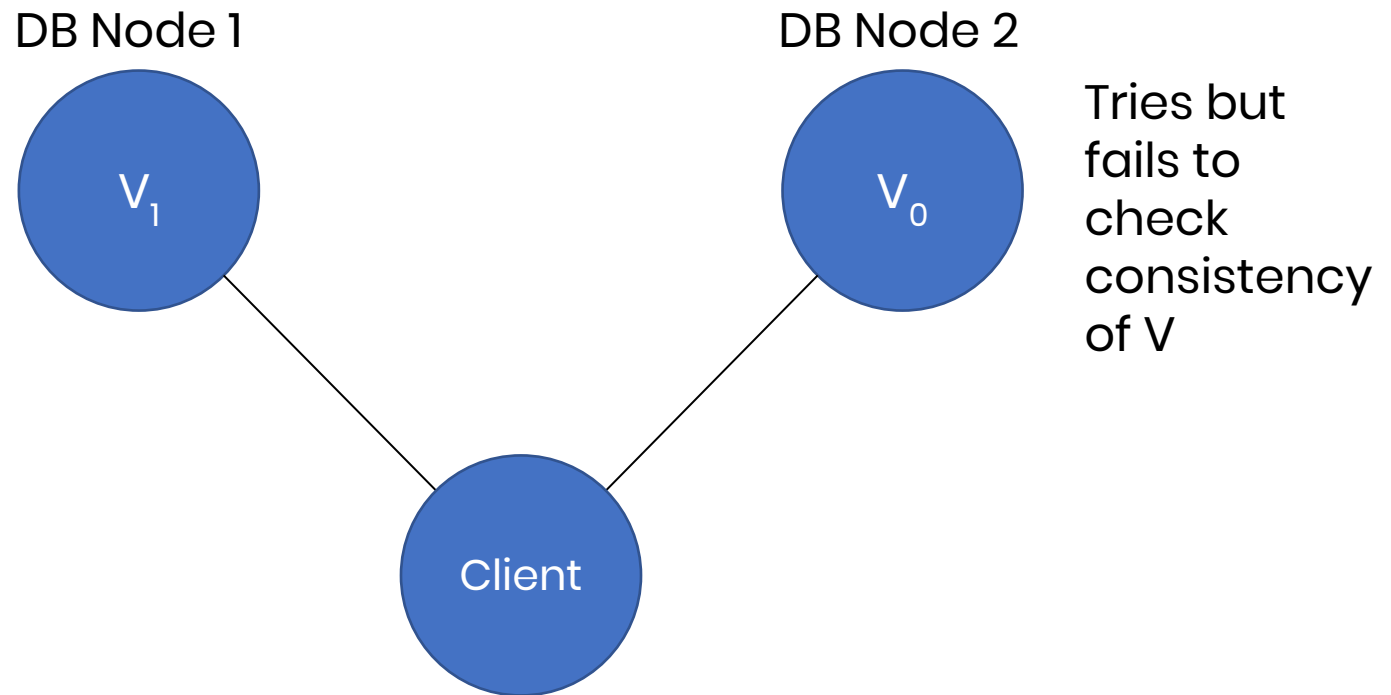
RDBMS vs NoSQL Systems

Partition tolerance + **Availability**



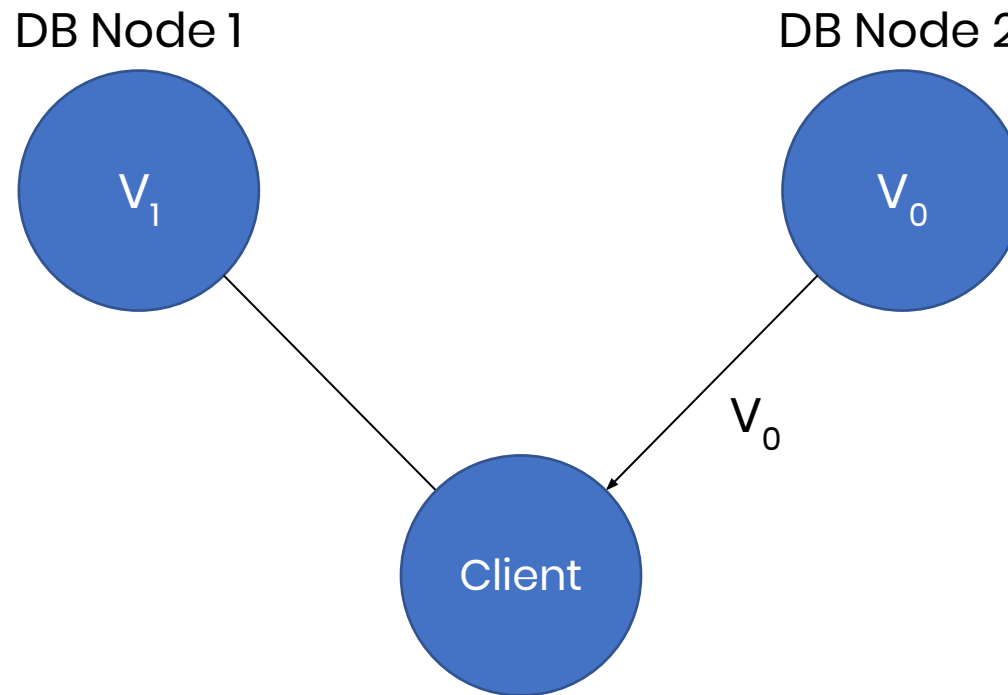
RDBMS vs NoSQL Systems

Partition tolerance + **Availability**



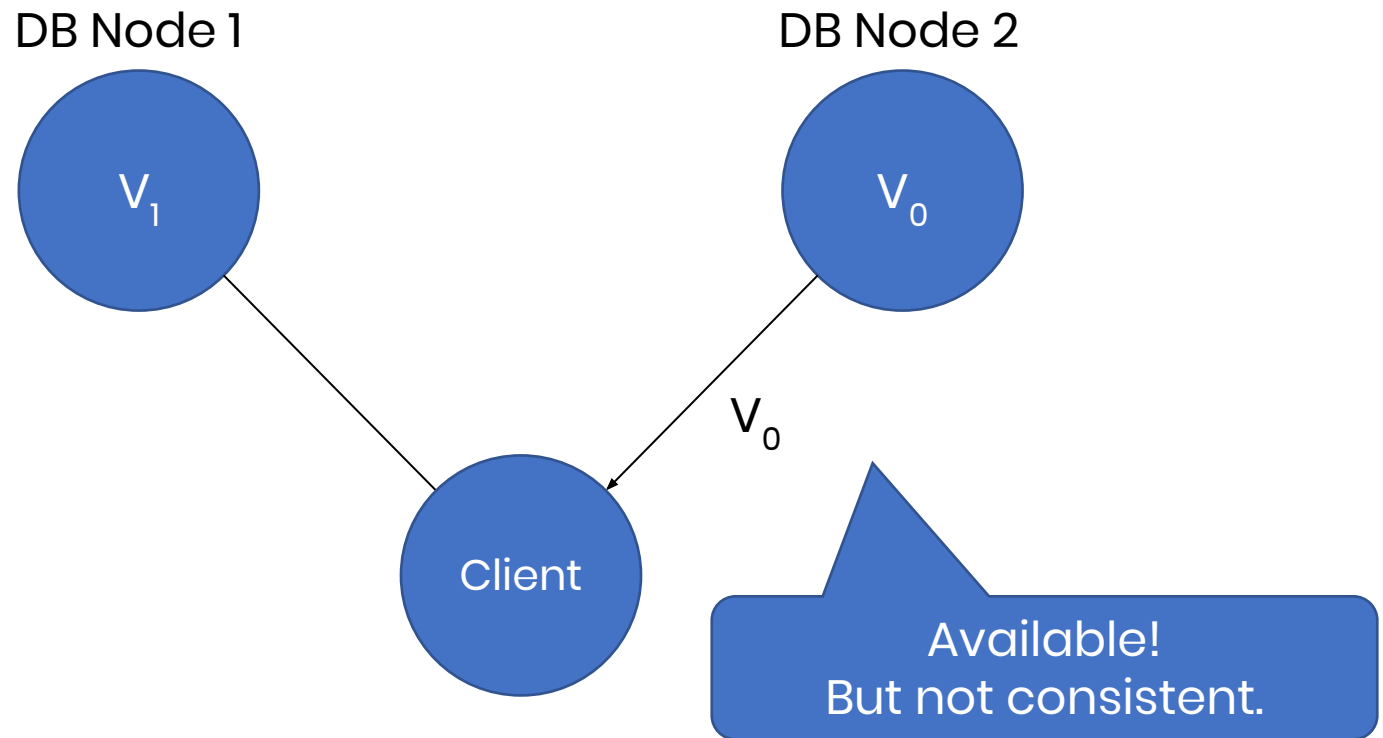
RDBMS vs NoSQL Systems

Partition tolerance + **Availability**



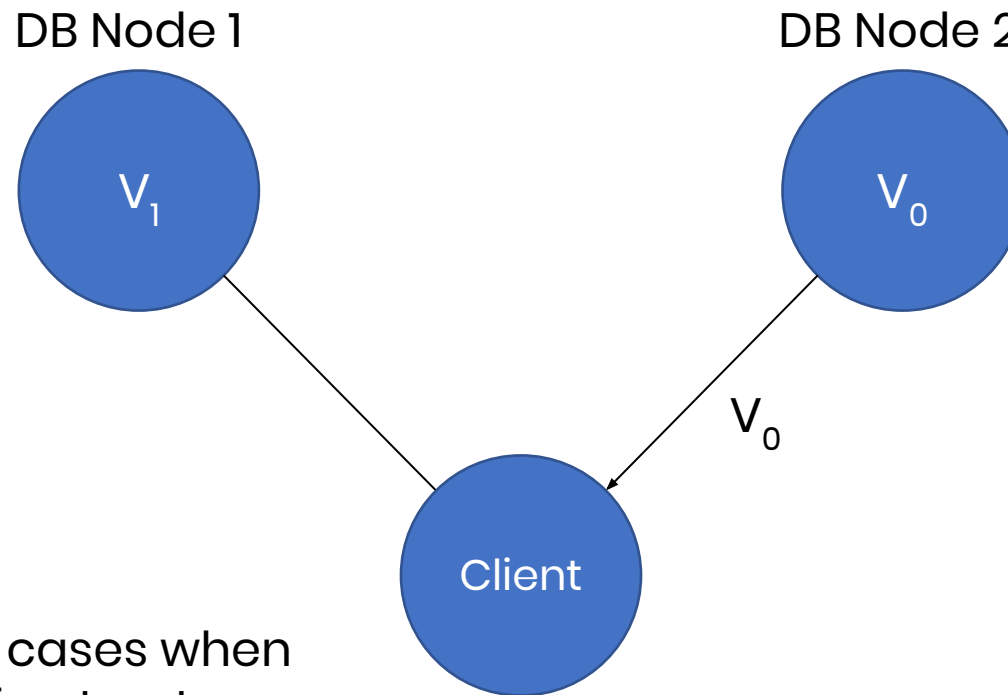
RDBMS vs NoSQL Systems

Partition tolerance + **Availability**



RDBMS vs NoSQL Systems

Partition tolerance + **Availability**



IMPORTANT:

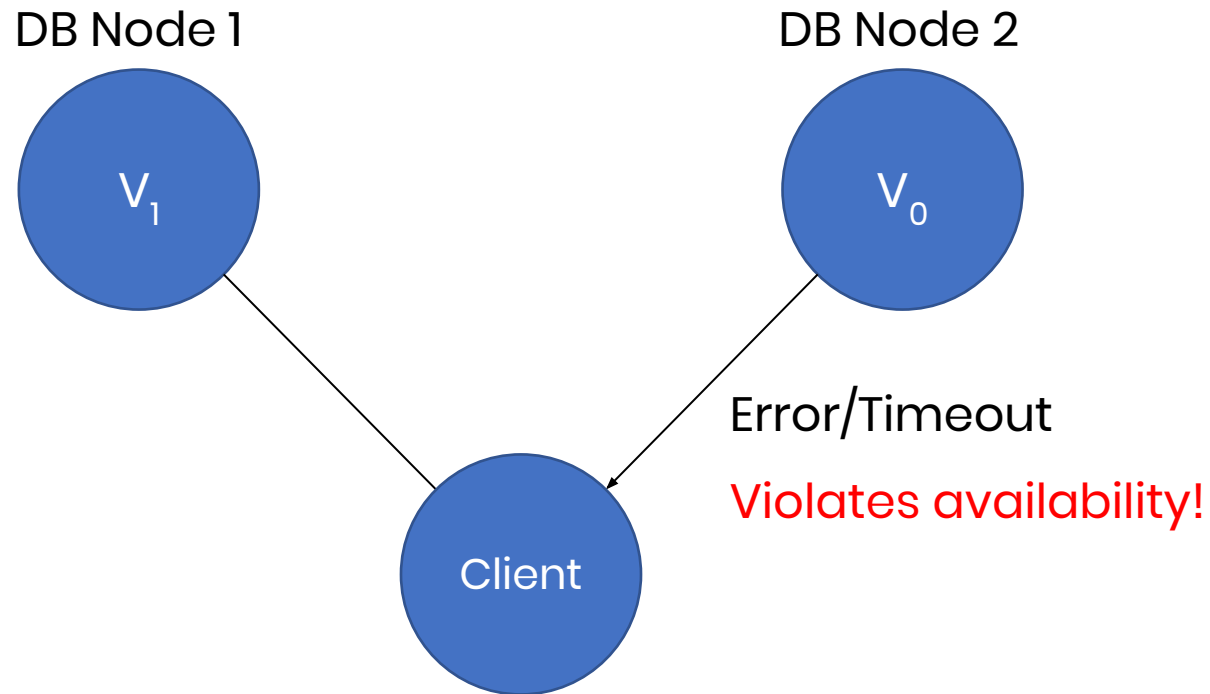
These are only cases when the network infrastructure goes down. Usually, nodes should be able to check on other nodes.

Proof of CAP Theorem

- [2002 original paper \(S. Gilbert & N. Lynch\)](#)
- [More digestible blog post \(M. Whittaker\)](#)
- Proof by contradiction: Assume we had a system that guaranteed availability, consistency, and partition tolerance...

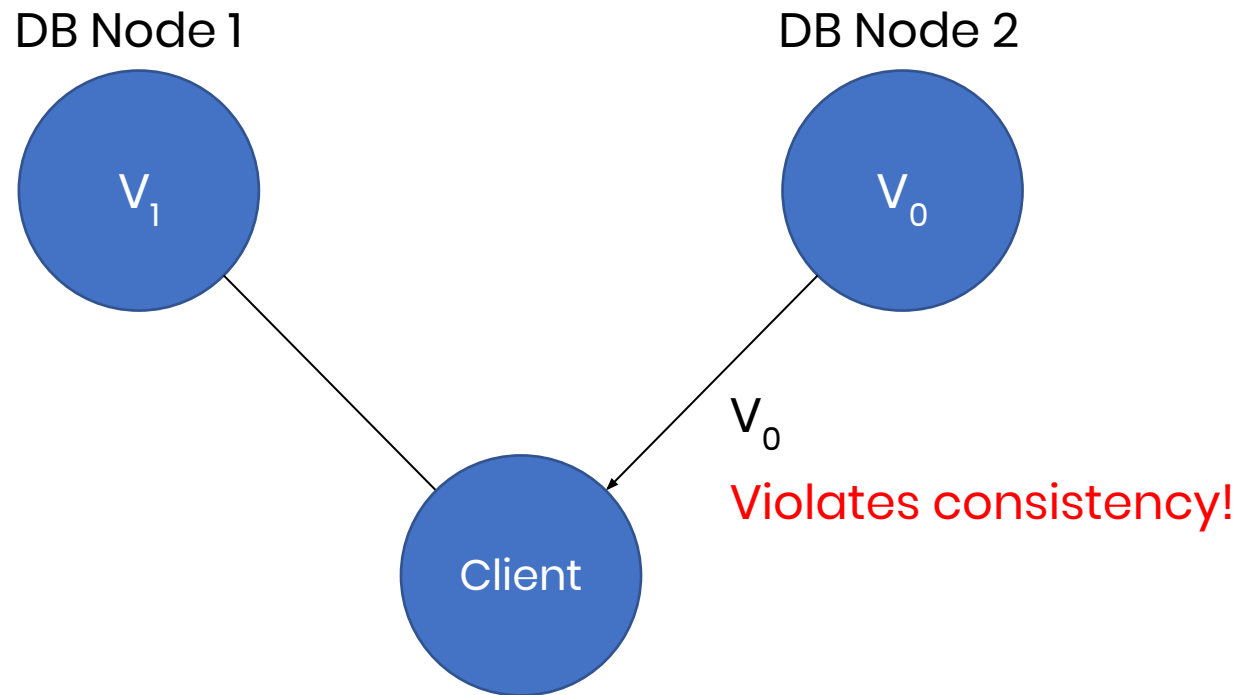
Proof of CAP Theorem

Partition tolerance + Consistency
+ Availability?



Proof of CAP Theorem

Partition tolerance + Availability
+ Consistency?





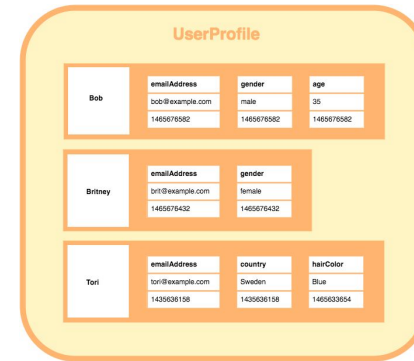
- RDBMSs are *intended* to be highly consistent
 - Boost availability by sacrificing some consistency
- NoSQL systems are *intended* to be highly available
 - Boost consistency by sacrificing some availability
- Most applications OK with some compromise
 - “Return most of data most of the time”
 - DBMS choice has many factors
 - Consistency/Availability requirements
 - Scalability
 - Usability
 - OLAP/analysis requirements
 - ...

NoSQL Data Models

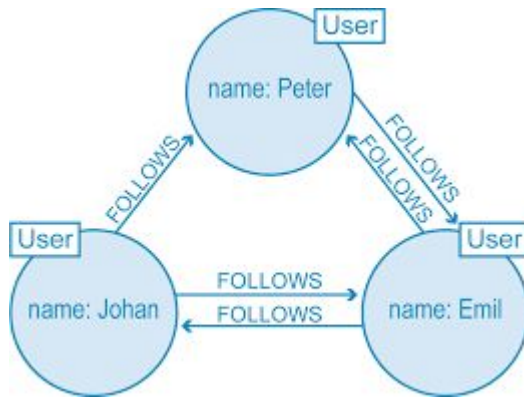
Key-Value Database

Key	Value
K1	AAA,BBB,CCC
K2	AAA,BBB
K3	AAA,DDD
K4	AAA,2,01/01/2015
K5	3,ZZZ,5623

Wide-Column Store (Extensible Record Store)



Graph Database



Document Store

XML

```
<empinfo>
  <employees>
    <employee>
      <name>James Kirk</name>
      <age>40</age>
    </employee>
    <employee>
      <name>Jean-Luc Picard</name>
      <age>45</age>
    </employee>
    <employee>
      <name>Wesley Crusher</name>
      <age>27</age>
    </employee>
  </employees>
</empinfo>
```

JSON

```
{ "empinfo" :
  {
    "employees" : [
      {
        "name" : "James Kirk",
        "age" : 40,
      },
      {
        "name" : "Jean-Luc Picard",
        "age" : 45,
      },
      {
        "name" : "Wesley Crusher",
        "age" : 27,
      }
    ]
  }
}
```

Coming up

- Deeper look at using NoSQL stores
- First we'll look at key-value stores
- Then we'll look at semi-structured data
- Then the query language SQL++ for AsterixDB